

Introductory Applied Machine Learning

Generalization, Overfitting, Evaluation

Victor Lavrenko and Nigel Goddard
School of Informatics


Generalization

- Training data: $\{x_i, y_i\}$
 - examples that we used to train our predictor
 - e.g. all emails that our users labelled ham / spam
- Future data: $\{x_i, ?\}$
 - examples that our classifier has never seen before
 - e.g. emails that will arrive tomorrow
- Want to do well on future data, not training
 - not very useful: we already know y_i
 - easy to be perfect on training data (DT, kNN, kernels)
 - does not mean you will do well on future data
 - can over-fit to idiosyncrasies of our training data

Under- and Over-fitting

- Over-fitting:

- predictor too complex (flexible)
 - fits “noise” in the training data
 - patterns that will not re-appear
- predictor F over-fits the data if:
 - we can find another predictor F'
 - which makes more mistakes on training data: $E_{train}(F') > E_{train}(F)$
 - but fewer mistakes on unseen future data : $E_{gen}(F') < E_{gen}(F)$




- Under-fitting:



- predictor too simplistic (too rigid)
- not powerful enough to capture salient patterns in data
- can find another predictor F' with smaller E_{train} and E_{gen}

Under- and Over-fitting examples

Regression:






predictor too inflexible:
cannot capture pattern




predictor too flexible:
fits noise in the data

Classification:



Flexible vs. inflexible predictors

- Each dataset needs different level of “flexibility”
 - depends on task complexity + available data
 - want a “knob” to get rigid / flexible predictors
- Most learning algorithms have such knobs:
 - regression: order of the polynomial
 - NB: number of attributes, limits on σ^2 , ε
 - DT: #nodes in the tree / pruning confidence
 - kNN: number of nearest neighbors
 - SVM: kernel type, cost parameter
- Tune to minimize **generalization error**



Training vs. Generalization Error

- Training error:

$$E_{train} = \frac{1}{n} \sum_{i=1}^n \overbrace{\text{error}(f_D(\mathbf{x}_i), y_i)}^{\text{same? different by how much?}}$$

training examples
value we predicted
true value

- Generalization error:

- how well we will do on future data
- don't know what future data x_i will be
- don't know what labels y_i it will have
- but know the “range” of all possible $\{x, y\}$
 - x : all possible 20x20 black/white bitmaps
 - y : $\{0, 1, \dots, 9\}$ (digits)

Usually

$$E_{train} \leq E_{gen}$$

Can never compute
generalisation error

$$E_{gen} = \int \underbrace{\text{error}(f_D(\mathbf{x}), y)}_{\text{error as before}} \underbrace{p(y, \mathbf{x})}_{\text{how often we expect to see such } x \text{ and } y} d\mathbf{x}$$

over all possible x, y


Estimating Generalization Error

- Testing error:
$$E_{test} = \frac{1}{n} \sum_{i=1}^n \text{error}(f_D(\mathbf{x}_i), y_i)$$
 over testing set
 - set aside part of training data (testing set)
 - learn a predictor without using any of this data
 - predict values for testing set, compute error
 - gives an **estimate** of true generalization error
 - if testing set is **unbiased** sample from $p(x,y)$: $\lim_{n \rightarrow \infty} E_{test} = E_{gen}$
 - how close? depends on n
- Ex: binary classification, 100 instances
 - assume: 75 classified correctly, 25 incorrectly
 - $E_{test} = 0.25$, E_{gen} around 0.25, but how close?

Confidence Interval for Future Error

- What range of errors can we expect for **future** test sets?
 - $E_{test} \pm \Delta E$ such that 95% of future test sets fall within that interval
- E_{test} is an unbiased estimate of $E = \text{true error rate}$
 - $E = \text{probability our system will misclassify a random instance}$
 - take a random set of n instances, how many misclassified?
 - flip E -biased coin n times, how many heads will we get?
 - Binomial distribution with mean $= n E$, variance $= n E (1-E)$
 - $E_{future} = \text{\#misclassified} / n$, \sim Gaussian, mean E , variance $= E(1-E) / n$
 - 2/3 future test sets will have error in $E \pm \sqrt{E(1-E)/n}$
- $p\%$ confidence interval for future error:
 - for $n=100$ examples, $p=0.95$ and $E = 0.25$
 - $\sigma = \sqrt{(0.25 \cdot 0.75 / 100)} = .043$
 - $CI = 0.25 \pm 1.96 \cdot \sigma = \mathbf{0.25 \pm 0.08}$
 - $n=100, p=0.99 \rightarrow CI = \mathbf{0.25 \pm 0.11}$
 - $n=10000, p=0.95 \rightarrow CI = \mathbf{0.25 \pm 0.008}$

our test set is one such set




Training, Validation, Testing sets


- Training set: construct classifier
 - NB: count frequencies, DT: pick attributes to split on
- Validation set: pick algorithm + knob settings
 - pick best-performing algorithm (NB vs. DT vs. ...)
 - fine-tune knobs (tree depth, k in kNN, c in SVM ...)
- Testing set: estimate future error rate
 - never report best of many runs
 - run only once, or report results of every run
- Split **randomly** to avoid bias

Cross-validation


- Conflicting priorities when splitting the dataset
 - estimate future error as accurately as possible
 - large testing set: big n_{test} \rightarrow tight confidence interval
 - learn classifier as accurately as possible
 - large training set: big n_{train} \rightarrow better estimates
 - training and testing cannot overlap: $n_{\text{train}} + n_{\text{test}} = \text{const}$
- Idea: evaluate Train \rightarrow Test, then Test \rightarrow Train, average results
 - **every** point is both training and testing, never at the same time
 - reduces chances of getting an unusual (biased) testing set
 - 5-fold cross-validation
 - randomly split the data into 5 sets
 - test on each in turn (train on 4 others)
 - average the results over 5 folds
 - more common: 10-fold



Fold 1



Fold 2



Fold 3

Leave-one-out

- n-fold cross-validation (n = total number of instances)
 - predict each instance, training on all $(n-1)$ other instances
- Pros and cons:
 - best possible classifier learned: $n-1$ training examples
 - high computational cost: re-learn everything n times
 - not an issue for instance-based methods like kNN
 - there are tricks to make such learning faster
 - classes not balanced in training / testing sets
 - random data, 2 equi-probable classes \rightarrow wrong 100% of the time
 - testing balance: {1 of A, 0 of B} vs. training: { $n/2$ of B, $n/2-1$ of A}
 - duplicated data \rightarrow nothing can beat 1NN (0% error)
 - wouldn't happen with 10-fold cross-validation

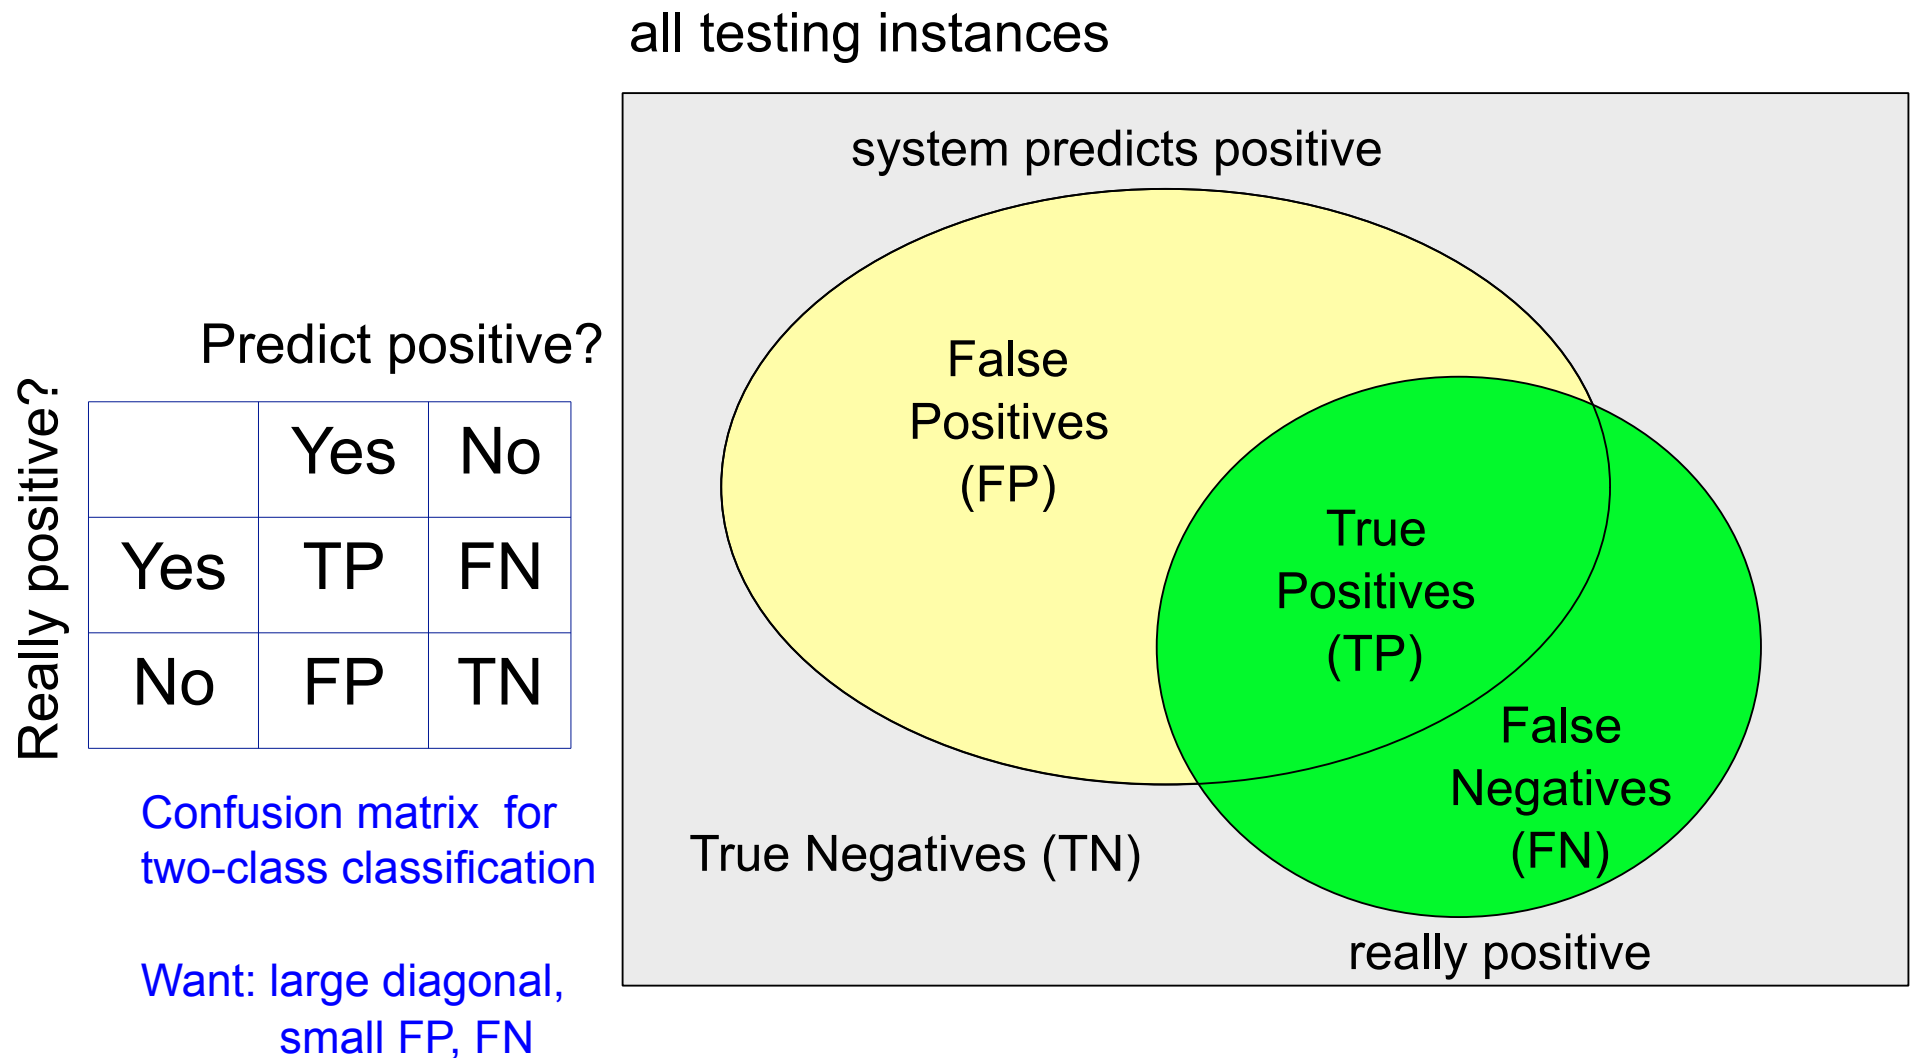
Stratification

- Problems with leave-one-out:
 - training / testing sets have classes in different proportions
 - not limited to leave-one-out:
 - K-fold cross-validation: random splits → imbalance
- Stratification
 - keep class labels balanced across training / testing sets
 - simple way to guard against unlucky splits
 - recipe:
 - randomly split each class into K parts
 - assemble i^{th} part from all classes to make the i^{th} fold


Evaluation measures

- Are we doing well? Is system A better than B?
- A measure of how (in)accurate a system is on a task
 - in many cases Error (Accuracy / PC) is not the best measure
 - using the appropriate measure will help select best algorithm
- Classification
 - how often we classify something right / wrong
- Regression
 - how close are we to what we're trying to predict
- Unsupervised
 - how well do we describe our data
 - in general – really hard

Classification measures: basics




Classification Error




- Classification error = $\frac{\text{errors}}{\text{total}} = \frac{FP + FN}{TP + TN + FP + FN}$
- Accuracy = $(1 - \text{error}) = \frac{\text{correct}}{\text{total}} = \frac{TP + TN}{TP + TN + FP + FN}$
- Basic measure of “goodness” of a classifier
- Problem: cannot handle unbalanced classes
 - ex1: predict whether an earthquake is about to happen
 - happen very rarely, very good accuracy if always predict “No”
 - solution: make FNs much more “costly” than FPs
 - ex2: web search: decide if a webpage is relevant to user
 - 99.9999% of pages not relevant to any query → retrieve nothing
 - solution: use measures that don’t involve TN (recall / precision)

Accuracy and un-balanced classes

- You're predicting Nobel prize (+) vs. not (•)
- Human would prefer classifier **A**.
- Accuracy will prefer classifier **B** (fewer errors)
- Accuracy poor metric here



Misses and False Alarms



- False Alarm rate = False Positive rate = $FP / (FP + TN)$
 - % of negatives we misclassified as positive
- Miss rate = False Negative rate = $FN / (TP + FN)$
 - % of positives we misclassified as negative
- Recall = True Positive rate = $TP / (TP + FN)$
 - % of positives we classified correctly (1 – Miss rate)
- Precision = $TP / (TP + FP)$
 - % positive out of what we predicted was positive
- Meaningless to report just one of these
 - trivial to get 100% recall or 0% false alarm
 - typical: recall/precision or Miss / FA rate or TP/FP rate

Evaluation (recap)

- Predicting class C (e.g. spam)

- classifier can make two types of mistakes:

- FP: false positives – non-spam emails mistakenly classified as spam
- FN: false negatives – spam emails mistakenly classified as non-spam
- TP/TN: true positives/negatives – correctly classified spam/non-spam

- common error/accuracy measures:

- Classification Error: $\frac{\text{errors}}{\text{total}} = \frac{FP + FN}{TP + TN + FP + FN}$

- Accuracy = 1-Error: $\frac{\text{correct}}{\text{total}} = \frac{TP + TN}{TP + TN + FP + FN}$

meaningless
if classes
imbalanced

- False Alarm = False Positive rate = $FP / (FP + TN)$

- Miss = False Negative rate = $FN / (TP + FN)$

- Recall = True Positive rate = $TP / (TP + FN)$

- Precision = $TP / (TP + FP)$

always report
in pairs, e.g.:
Miss / FA or
Recall / Prec.

		Predicted C?	
Really C?		Yes	No
	Yes	TP	FN
	No	FP	TN

Utility and Cost


- Sometimes need a single-number evaluation measure
 - optimizing the learner (automatically), competitive evaluation
 - may know costs of different errors, e.g. earthquakes:
 - false positive: cost of preventive measures (evacuation, lost profit)
 - false negative: cost of recovery (reconstruction, liability)
- Detection cost: weighted average of FP, FN rates
 - $\text{Cost} = C_{\text{FP}} * \text{FP} + C_{\text{FN}} * \text{FN}$ [event detection]
- F-measure: harmonic mean of recall, precision
 - $F1 = 2 / (1 / \text{Recall} + 1 / \text{Precision})$ [Information Retrieval]
- Domain-specific measures:
 - e.g. observed profit/loss from +/- market prediction

Thresholds in Classification

- Two systems have the following performance:
 - A: True Positive = 50%, False Positive = 20%
 - B: True Positive = 100%, False Positive = 60%
- Which is better? (assume no-apriori utility)
 - very misleading question
 - A and B could be the same exact system
 - operating at different thresholds

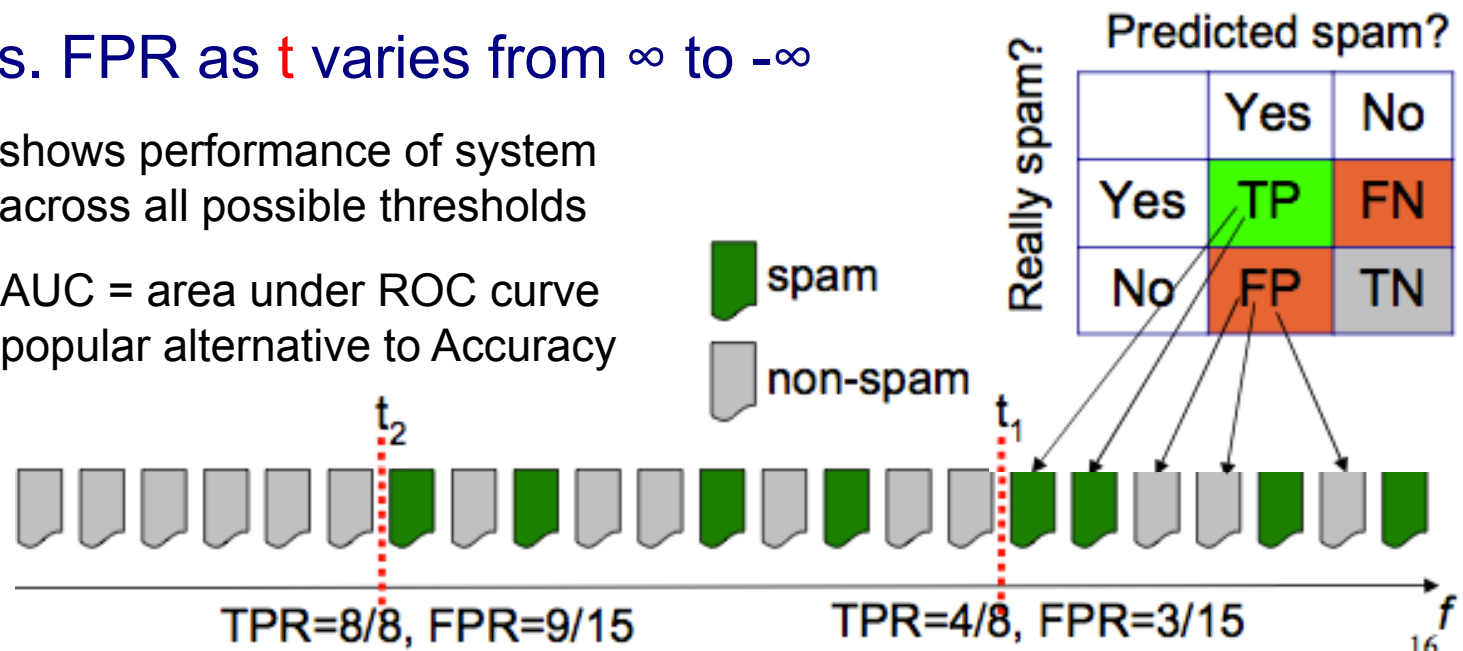
ROC curves

- Many algorithms compute “confidence” $f(x)$
 - threshold to get decision: spam if $f(x) > t$, non-spam if $f(x) \leq t$
 - Naïve Bayes: $P(\text{spam}|x) > 0.5$, Linear/Logistic/SVM: $w^T x > 0$, Decision Tree: $p_+/p_- > 1$
 - threshold t determines error rates
 - False Positive rate = $P(f(x) > t | \text{ham})$, True Positive rate = $P(f(x) > t | \text{spam})$
- Receiver Operating Characteristic (ROC):
 - plot TPR vs. FPR as t varies from ∞ to $-\infty$




shows performance of system across all possible thresholds

AUC = area under ROC curve
popular alternative to Accuracy




ROC convex hull

- System A: better at high thresholds (high-precision)
- System B: better at low thresholds (high-recall)
- System C: for each x : flip a p -coin, heads: $A(x)$, tails: $B(x)$
 - if x was really positive:
 - $P(\text{correct}) = p * P(A(x) > t_A \mid +) + (1-p) * P(B(x) > t_B \mid +)$
 - $TP_C = p TP_A + (1-p) TP_B$
 - if x was really negative:
 - $P(\text{error}) = p * P(A(x) > t_A \mid -) + (1-p) * P(B(x) > t_B \mid -)$
 - $FP_C = p FP_A + (1-p) FP_B$
 - may be better than either A or B
 - example: Netflix challenge



Evaluating regression

- Classification:
 - count how often we are wrong
- Regression:
 - predict numbers y_i from inputs x_i
 - always wrong, but by how much?
 - distance between predicted & true values
 - (root) mean squared error:
 - popular, well-understood, nicely differentiable
 - sensitive to single large errors (outliers)
 - mean absolute error:
 - less sensitive to outliers
 - correlation coefficient
 - insensitive to mean & scale



$$\sqrt{\frac{1}{n} \sum_{i=1}^n \underbrace{f(x_i)}_{\text{predicted}} - \underbrace{y_i}_{\text{true}})^2}$$

testing set


$$\frac{1}{n} \sum_{i=1}^n |f(x_i) - y_i|$$

$$\frac{n \sum_{i=1}^n (f(x_i) - \mu_f)(y_i - \mu_y)}{\sqrt{\sum_{i=1}^n (f(x_i) - \mu_f)^2 \cdot \sum_{i=1}^n (y_i - \mu_y)^2}}$$

Mean Squared Error

- Average (squared) deviation from truth $\sqrt{\frac{1}{n} \sum_{i=1}^n (f(x_i) - y_i)^2}$
- Very sensitive to outliers
 - 99 exact, 1 off by \$10
 - all 100 wrong by \$1

} same MSE \Rightarrow large effect on models
- Sensitive to mean / scale
 - $\mu_y = 1/n \sum_i y_i$... good baseline
- Relative squared error (Weka)



$$\sqrt{\frac{\sum_{i=1}^n (f(x_i) - y_i)^2}{\sum_{i=1}^n (\mu_y - y_i)^2}}$$


} MSE of predictor

} MSE when using the mean as a predictor

Mean Absolute Error

- Mean Absolute Error (MAE):

- less sensitive to outliers
- many small errors = one large error
- best 0th order baseline: $\text{median}\{y_i\}$
 - not the mean as for MSE

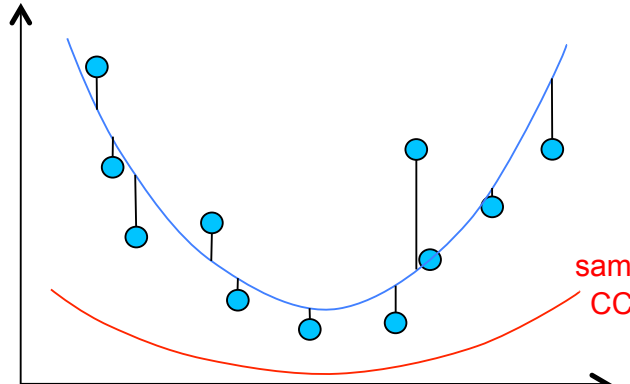


- Median Absolute Deviation (MAD): $\text{med}\{|f(x_i) - y_i|\}$

- robust, completely ignores outliers
 - can define similar squared error: $\text{median}\{(f(x_i) - y_i)^2\}$
 - difficult to work with (can't take derivatives)
- Sensitive to mean, scale

Correlation Coefficient

- Completely insensitive to mean / scale:

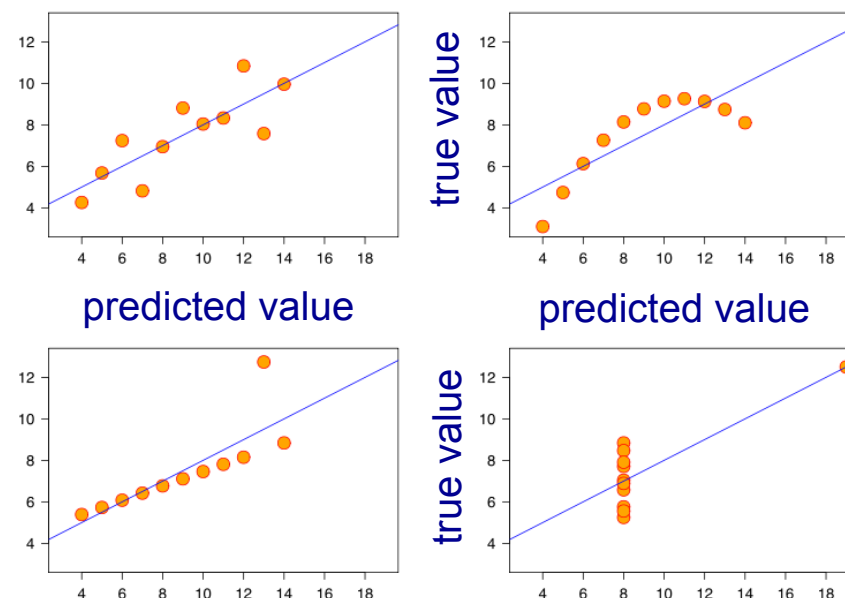
$$\frac{n \sum_{i=1}^n (f(x_i) - \mu_f)(y_i - \mu_y)}{\sqrt{\sum_{i=1}^n (f(x_i) - \mu_f)^2 \cdot \sum_{i=1}^n (y_i - \mu_y)^2}} = n \sum_{i=1}^n \underbrace{\frac{f(x_i) - \mu_f}{\sigma_f}}_{\text{prediction relative to mean}} \cdot \underbrace{\frac{y_i - \mu_y}{\sigma_y}}_{\text{truth relative to mean}}$$


- Intuition: did you capture the relative ordering?

- output larger $f(x_i)$ for larger y_i
- output smaller $f(x_i)$ for smaller y_i
- useful for ranking tasks:
 - e.g. recommend a movie to a user

- Important to visualize data

- same CC for 4 predictors →




Summary

- Training vs. generalization error
 - under-fitting and over-fitting
- Estimate how well your system is doing its job
 - how does it compare to other approaches?
 - what will be the error rate on future data?
- Training and testing
 - cross-validation, leave-one-out, stratification, significance
- Evaluation measures
 - accuracy, miss / false alarm rates, detection cost
 - ROC curves
 - regression: (root) mean squared/absolute error, correlation

Evaluating unsupervised methods

- Generally hard and subjective
 - broad aim: did we capture the structure of the dataset?
 - if possible: does it help us do some (supervised) task
- Dimensionality reduction
 - distance between data in original & reduced space
- Mixture models
 - do we assign high probability to the training data?
- Clustering
 - did we “discover” the latent sub-populations?

Significance tests

- Often need to compare two systems: A, B
 - perform cross-validation: errors $e_{A,1} \dots e_{A,K}, e_{B,1} \dots e_{B,K}$
 - average errors: $e_A < e_B$ 
 - does this mean that A better than B?
 - look at the variance of errors
- Significance: could the difference be due to chance?
 - analogy: 3 coin flips, always large difference, pure chance
 - null hypothesis H_0 :
 - $e_{A,1} \dots e_{A,K}, e_{B,1} \dots e_{B,K}$ are random samples from the same population
 - want to show $P(H_0)$ is very small \rightarrow reject H_0 as improbable
 - let $d_i = e_{A,i} - e_{B,i}$ $t = \frac{\sum_i d_i}{\sqrt{\sum_i (d_i - \mu)^2}} \sim$ Student's t distribution
 - caution: d_i must be independent (no overlap in data)