

IAML: K-means Clustering

Victor Lavrenko and Nigel Goddard

School of Informatics


Semester 1

Overview

- Clustering
- K-means algorithm
- Practical issues: local optimum, selecting K
- Evaluating clustering algorithms
- Application: image representation
- Reading:
 - Witten & Frank sections 4.8 and 6.6

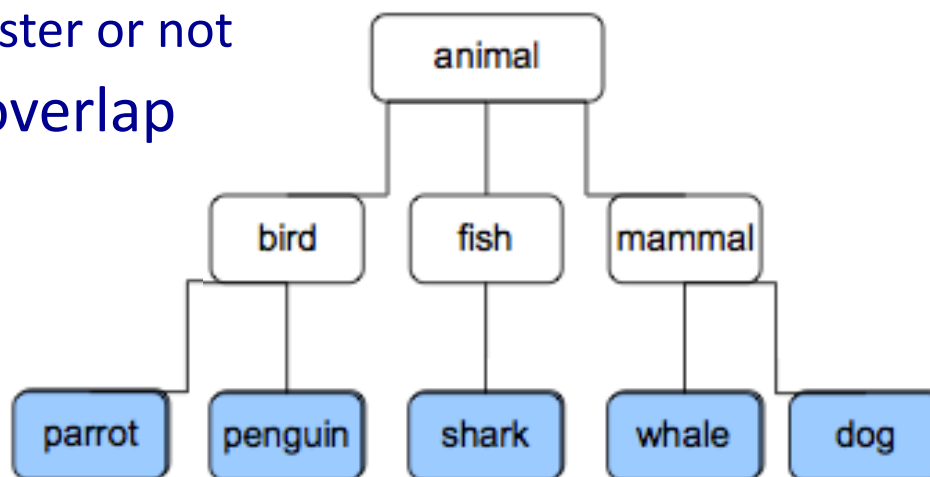
Clustering

- Discover the underlying structure of the data
 - unsupervised task, not predicting anything specific
- What sub-populations exist in the data?
 - how many are there?
 - what are their sizes?
 - do elements in a sub-population have any common properties?
 - are sub-populations cohesive?
 - can they be further split up?
 - are there outliers?




Types of clustering methods

- Goal:
 - monothetic: cluster members have some common property
 - e.g. all are males aged 20-35, or all have X% response to test B
 - polythetic: cluster members are similar to each other
 - distance between elements defines membership
- Overlap:
 - hard clustering: clusters do not overlap
 - element either belongs to a cluster or not
 - soft clustering: clusters may overlap
 - “strength of association”
between element and cluster
- Flat or hierarchical
 - set of groups vs. taxonomy



Methods we will cover

- K-D Trees (see k-NN lecture)
 - monothetic, hard boundaries, hierarchical
- K-means clustering
 - splits data into a specified number of populations
 - polythetic, hard boundaries, flat
- Gaussian mixtures (EM algorithm)
 - fits a mixture of K Gaussians to the data
 - polythetic, soft boundaries, flat
- Agglomerative clustering
 - creates an “ontology” of nested sub-populations
 - polythetic, hard boundaries, hierarchical



K-means clustering



- Produces hard, flat, polythetic clusters
 - data partitioned into K sub-populations (need to know K)
 - points in each sub-population similar to a “centroid”
 - centroid = attribute-value “representation” of a cluster
 - “prototypical” individual in a sub-population
- Uses:
 - discover classes in an unsupervised manner
 - e.g. cluster images of handwritten digits (with $K = 10$)
 - smoothness over space
 - in the same cluster → similar representations / class labels / ...
 - dimensionality reduction: clusters = “latent factors”
 - replace representation of each data point with its cluster number
 - assumes all pertinent qualities reflected in cluster membership
 - related to basis / kernels in linear classifiers

K-means clustering algorithm

- Input: K , set of points $x_1 \dots x_n$
- Place centroids $c_1 \dots c_K$ at random locations

- Repeat until convergence:

- for each point x_i :

- find nearest centroid c_j $\arg \min_j D(x_i, c_j)$
 - assign the point x_i to cluster j

distance (e.g. Euclidian) between
instance x_i and cluster center c_j

- for each cluster $j = 1 \dots K$:

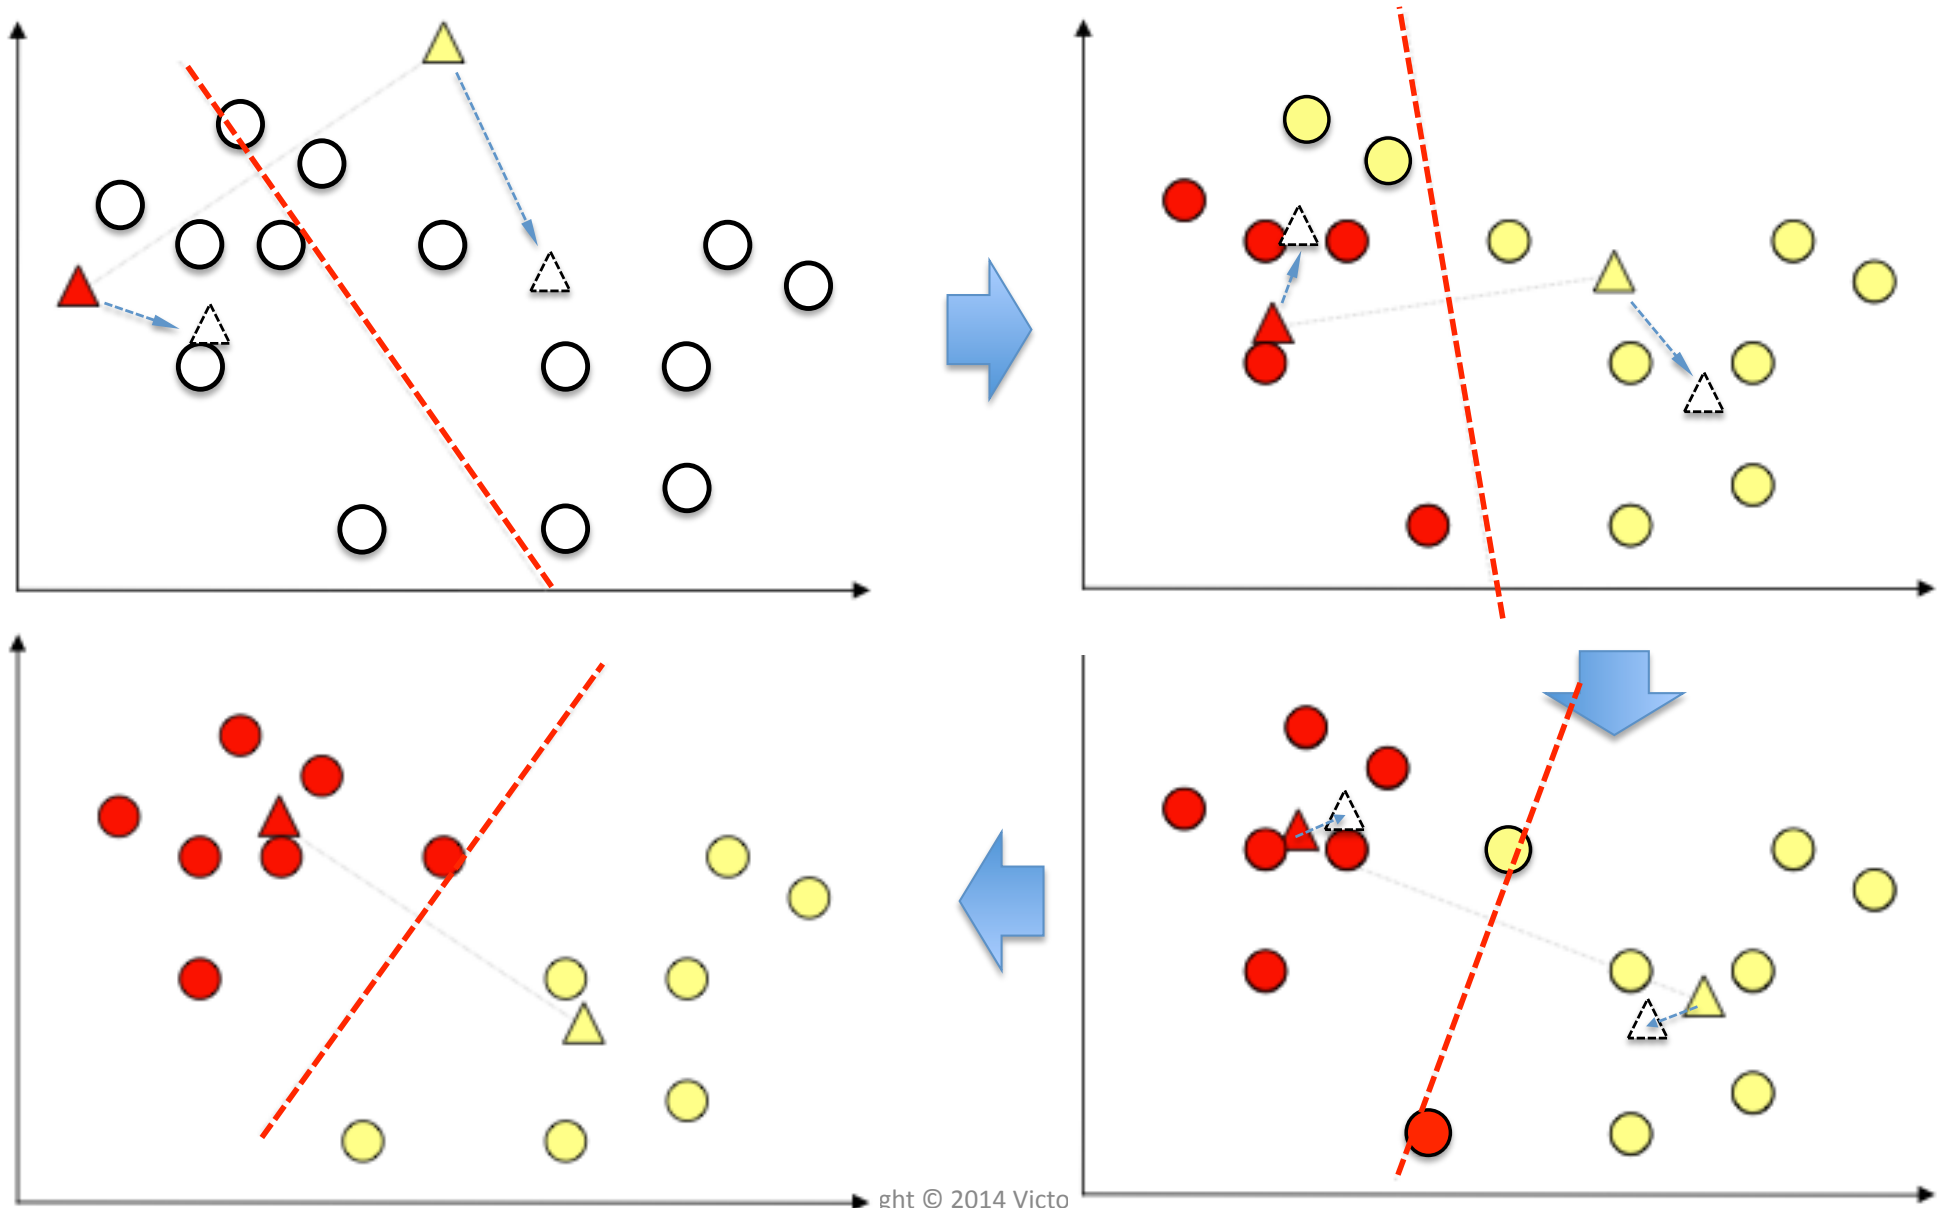
- new centroid c_j = mean of all points x_i
assigned to cluster j in previous step

$$c_j(a) = \frac{1}{n_{jx_i \rightarrow c_j}} \sum x_i(a) \quad \text{for } a = 1 \dots d$$

- Stop when none of the cluster assignments change

$O(\text{\#iterations} * \text{\#clusters} * \text{\#instances} * \text{\#dimensions})$

K-means clustering example



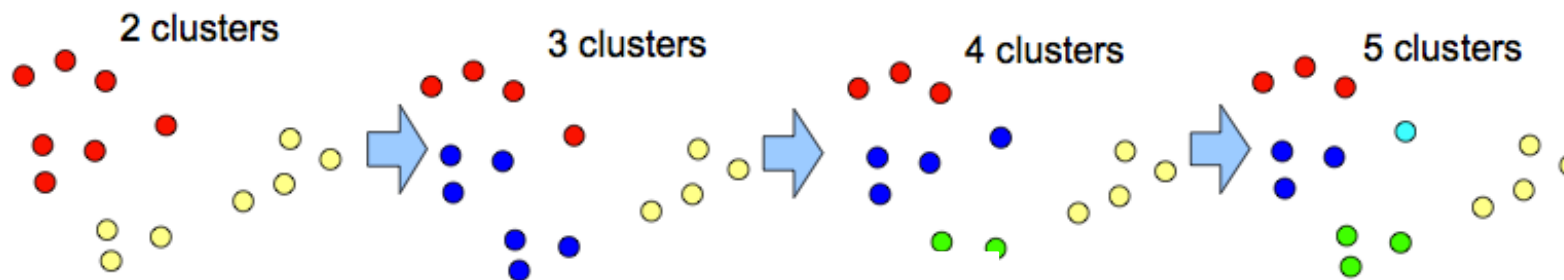

K-means properties

- Minimizes aggregate intra-cluster distance $\sum_j \sum_{x_i \rightarrow c_j} D(c_j, x_i)^2$
 - total squared distance from point to centre of its cluster
 - same as variance if Euclidian distance is used
- Converges to a local minimum
 - different starting points → very different results
 - run several times with random starting points
 - pick clustering that yields smallest aggregate distance
- Nearby points may not end up in the same cluster
 - the following clustering is a stable local minimum:



Optimal number of clusters

- How many clusters are there in your data?
 - class labels may suggest the value of K (e.g. digits 0..9)
 - optimize distance V: for $K = 2, 3, \dots$
 - run K-means, record distance
 - problem: V minimized when $K = n$
 - what if we use a validation set?
 - W&F: Minimum Description Length
 - total bits to encode K centroids + V
 - visually from scree plot:
 - point where “mountain” ends, “rubble” begins
 - elbow method: maximize 2nd derivative of V:
point where rate of decline changes the most




Evaluating Clustering Algorithms

- **Extrinsic** (helps us solve another problem)
 - represent images with cluster features
 - train different classifier for each sub-population
 - identify and eliminate outliers / corrupted points
- **Intrinsic** (useful in and of itself)
 - helps understand the makeup of our data (qualitative)
 - clusters correspond to classes (digits → 10 clusters)
 - align, evaluate as you would a normal classifier
 - compare to human judgments
 - can't ask humans to “cluster” a dataset manually
 - sample pairs x_i, x_j ask humans if they “match”

did your
classifier
improve?

Intrinsic Evaluation 1

- System produces clusters $C_1 C_2 \dots C_K$
- Reference clusters (classes) $R_1 R_2 \dots R_N$
- Align up $R_i \leftrightarrow C_j$, measure accuracy, F1, ...
 - many different ways to align:
 - Weka: $C_j \rightarrow R_i$ with max overlap
 - if many $C_j \rightarrow$ same R_i :
 - re-assign in a greedy manner
 - non-greedy: $K!/(N-K)!$ ways (very slow)
 - can we have multiple $C_j \rightarrow$ same R_i ?
 - can we have multiple $R_i \rightarrow$ same C_j ?
 - can we have overlapping clusters?




true class \rightarrow	R2	R1	R3	
cluster \rightarrow C1	3	1	2	6
C2	0	0	1	1
C3	7	1	8	16
C4	2	0	1	3
	12	2	12	

Accuracy = $(3+0+8)/26$

Intrinsic Evaluation 2


- Sample pairs x_i, x_j
 - ask human if x_i, x_j should be in the same group
 - easy task (cognitively)
 - can't ask them to "cluster" dataset manually
- System produces clusters
- Count errors, compute accuracy, F1, etc
 - FN: **matching** pairs x_i, x_j that are in different clusters (**e,h**)
 - FP: **non-matching** pairs x_i, x_j that are in same cluster (**c,d**)
- Doesn't require a pairing strategy
- Can handle overlapping clusters (a bit tricky)
 - same pair can count as both TN and FP (**g,h = No**)
- Can generate pairs from classes

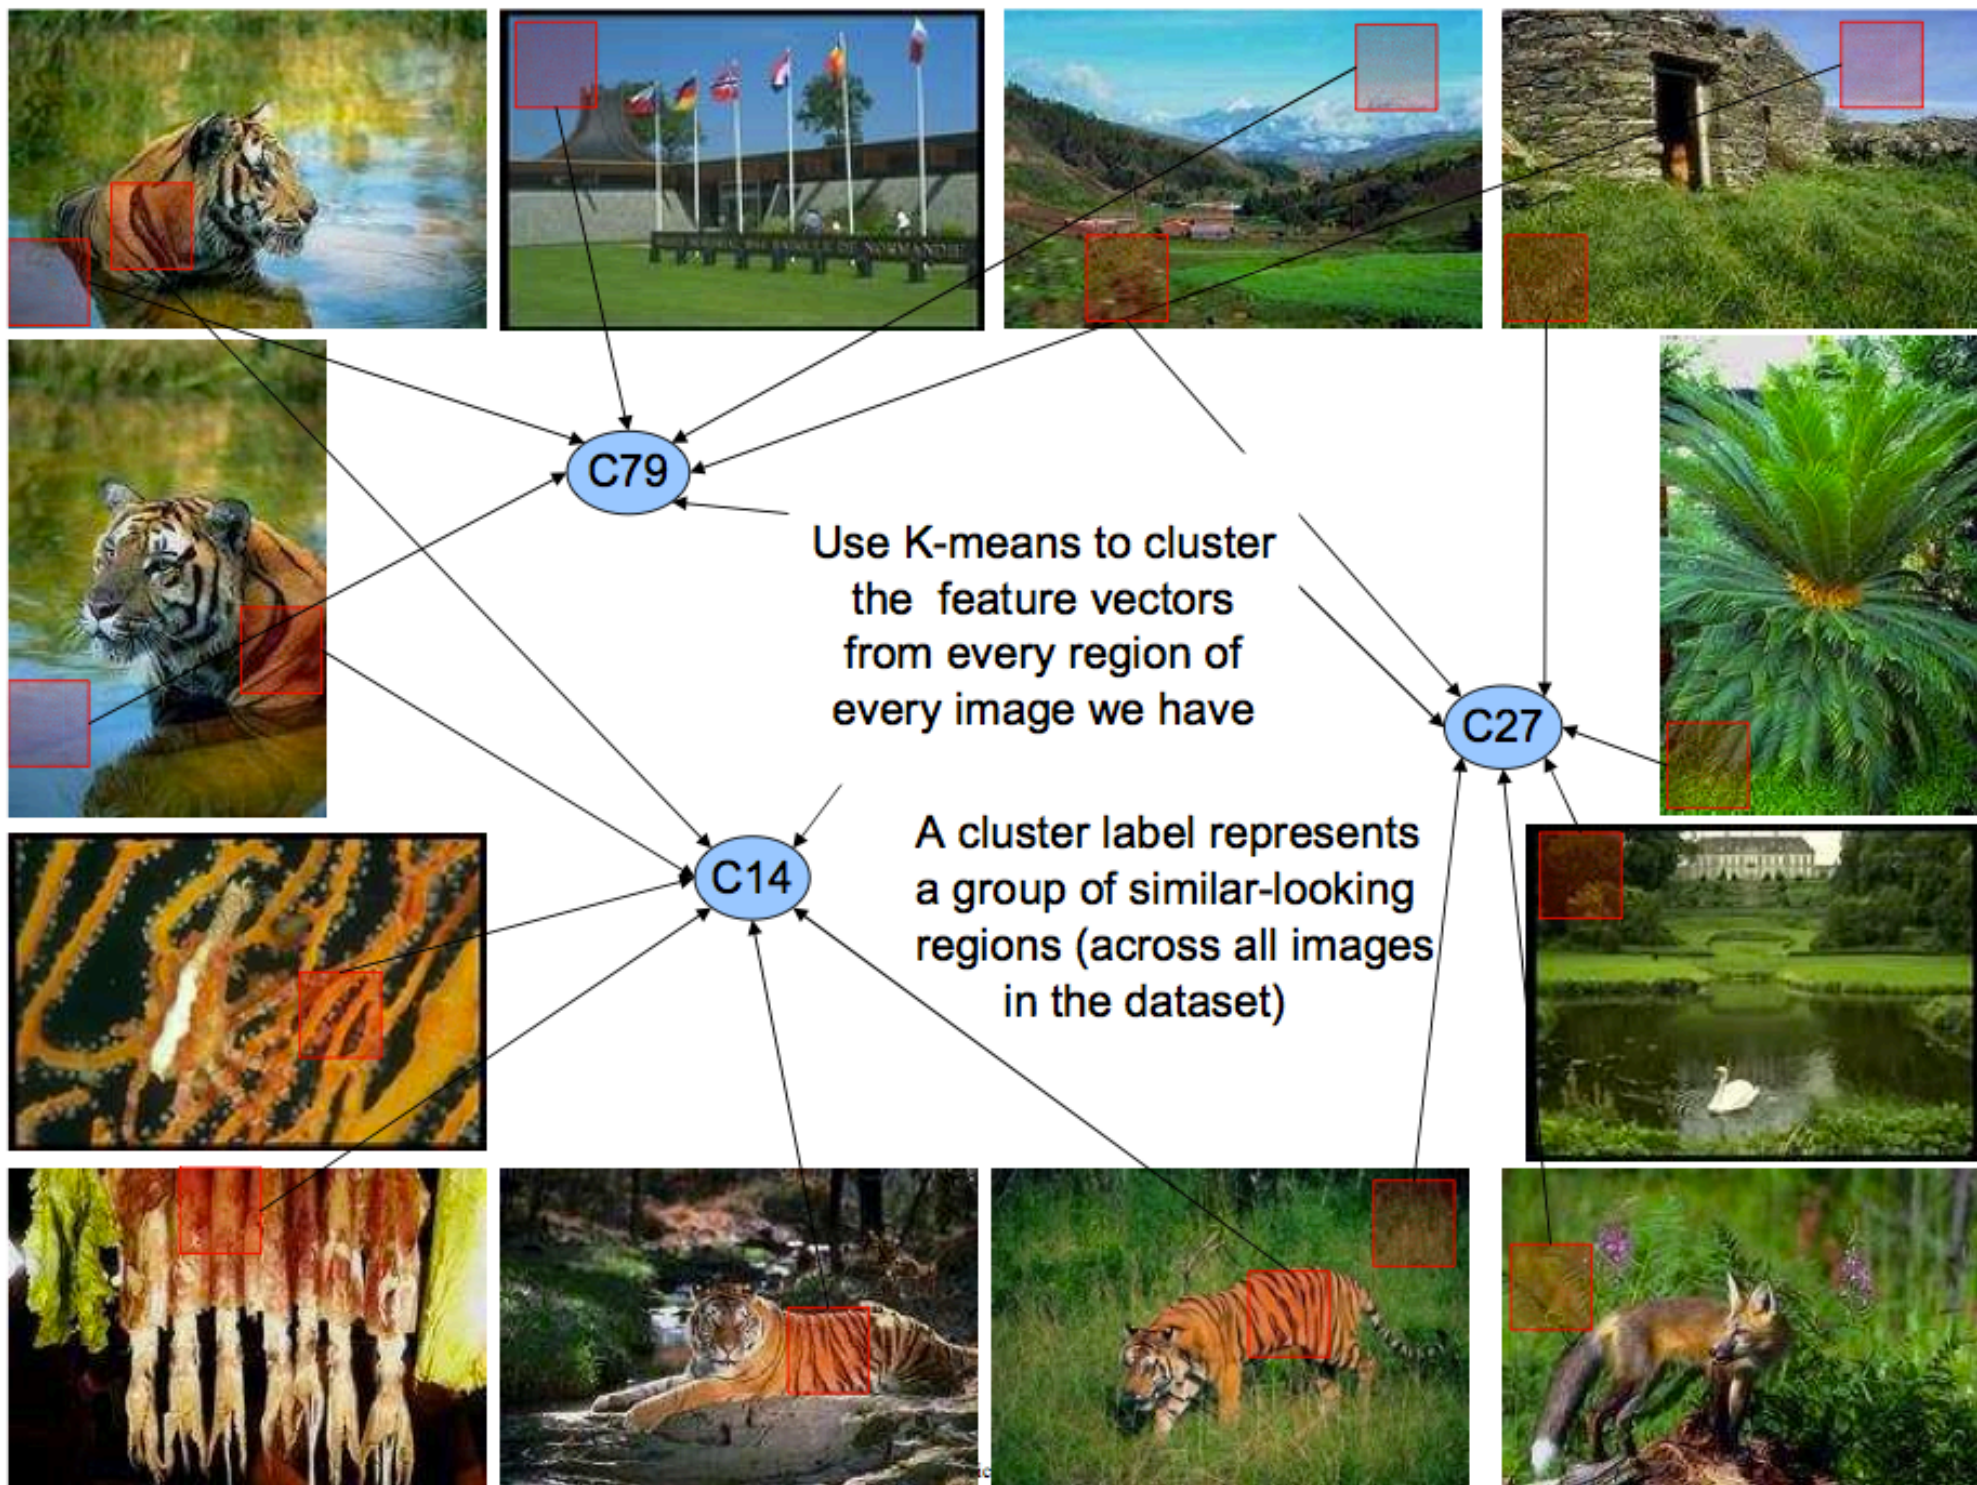
a,b = Yes
c,d = No
e,h = Yes
g,h = No



Application: image representation

- Goal: detect presence / absence of objects in image
- First step: represent images as attribute-value pairs
 - pixels as attributes: $10^3 \times 10^3 \times 10^3$ (conservative)
 - large and not very meaningful for learning
- bag-of-words would be nice
 - {"water", "grass", "tiger", "cat", "ripples"}
 - requires human annotation
- break image into a set of patches
 - patch = part of some object
- compute appearance features for each patch
 - relative position, distribution of colors, texture, edge orient.
- convert to a "word" that reflects appearance of a patch
 - similar-looking feature vectors → same word to represent them





K-means for image representation


- What K-means does:
 - groups all feature vectors from all images into K clusters
 - provides a cluster id for every patch in every image
 - represents the salient properties of the patch
 - similar-looking patches have the same id

- Represent patch with cluster id

- image = bag of cluster ids
 - one for each patch in the image
- K-dimensional representation:

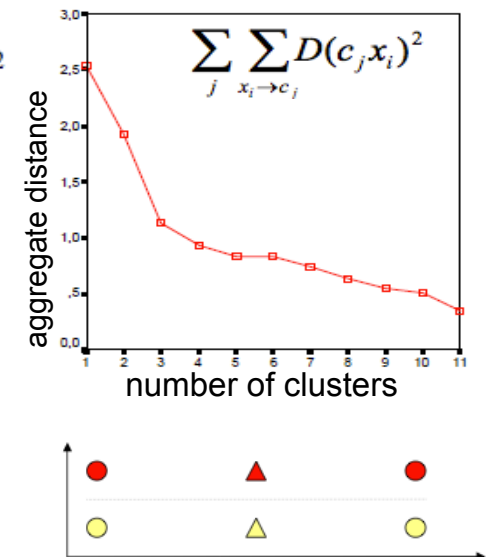
$\{ 4 \times \text{"C14"}, 7 \times \text{"C27"}, 24 \times \text{"C79"}, 0 \times \text{everything else} \}$

- similar to bag-of-words
- cluster ids sometimes called vis-terms or “visual words”



Summary

- Clustering: discover underlying sub-populations
- K-means
 - fast, iterative method: $O(i \cdot K \cdot n \cdot d)$
 - converges to a local minimum of $\sum_j \sum_{x_i \rightarrow c_j} D(c_j, x_i)^2$
 - run several times with different starting points
 - need to pick K: use scree plot
 - need to pick distance function (Euclidean)
 - nearby points may end up in diff. clusters
- Application: image representation
 - cluster image patches based on visual similarity
 - cluster numbers (vis-terms) becomes attributes
- Evaluation: intrinsic vs. extrinsic



Clustering: general structure

- Task: unsupervised / generative
 - group instances into K clusters
- Model structure
 - K cluster centroids (d -dimensional vectors)
- Score function
 - average distance from instance to cluster centre
- Optimization / search method
 - iteratively re-assign instances to clusters and update cluster centroids