



THE UNIVERSITY
of EDINBURGH

Text Technologies for Data Science

INFR11145

Indexing (2)

Instructor:
Walid Magdy

Lecture Objectives

- Learn more about indexing:
 - Structured documents
 - Extent index
 - Index compression
- Data structure
- Wild-char search and applications


** You are not asked to implement any of the content in this lecture, but you might think of using some for your course project 😊*

Structured Documents

- Document are not always flat:
 - Meta-data: title, author, time-stamp
 - Structure: headline, section, body
 - Tags: link, hashtag, mention
- How to deal with it?
 - Neglect!
 - Create separate index for each field
 - Use “extent index”

Extent Index

- Special “term” for each element/field/tag
 - Index all terms in a structured document as plain text
 - Terms in a given field/tag get special additional entry
 - Posting: spans of window related to a given field
 - Allows multiple overlapping spans of different types



D1: He likes to wink, he likes to drink

D2: He likes to drink, and drink, and drink


D3: The thing he likes to drink is ink

D4: The ink he likes to drink is pink

D5: He likes to wink, and drink pink ink

Using Extent

- Doc: 1 →
Headline: “*Information retrieval lecture*”
Text: “~~this is lecture 6 of the TTSD course on IR~~”
4 5 6 7 8
- Query → Headline: lecture


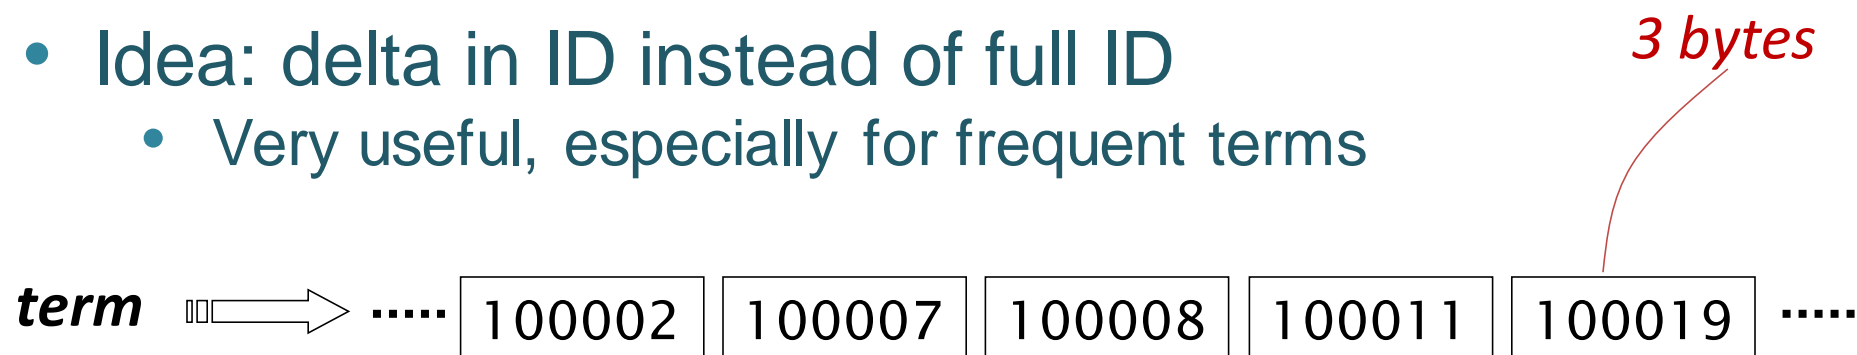


Index Compression

- Inverted indices are big
 - Large disk space → large I/O operations
- Index compression
 - Reduce space → less I/O
 - Allow more chunks of index to be cached in memory
- Large size goes to:
 - terms? document numbers?
 - Ideas:
 - Compress document numbers, how?

Delta Encoding


- Large collections → large sequence of doc IDs
- Large ID number → more bytes to store
 - 1 byte: 0 → 255
 - 2 bytes: 0 → 65,535
 - 4 bytes: 0 → 4.3 B
- Idea: delta in ID instead of full ID
 - Very useful, especially for frequent terms



v-byte Encoding

- Have different byte storage for each delta in index
 - Use fewer bits to encode
 - High bit in a byte \rightarrow 1/0 = terminate/continue
 - Remaining 7 bits \rightarrow binary number
 - Examples:
 - “6” \rightarrow 10000110
 - “127” \rightarrow 11111111
 - “128” \rightarrow ~~0~~00000001~~1~~00000000

- Real example sequence:



Index Compression

- There are more sophisticated compression algorithms:
 - Elias gamma code
- The more compression
 - Less storage
 - More processing
- In general
 - Less I/O + more processing > more I/O + no processing
“>” = faster
 - With new data structures, problem is less severe


Dictionary Data Structures

- The dictionary data structure stores the term vocabulary, document frequency, pointers to each postings list ...
- For small collections, load full dictionary in memory. In real-life, cannot load all index to memory!
 - Then what to load?
 - How to reach quickly?
 - What data structure to use for inverted index?


Hashes

- Each vocabulary term is hashed to an integer
- Pros
 - Lookup is faster than for a tree: $O(1)$
- Cons
 - No easy way to find minor variants:
 - judgment/judgement
 - No prefix search
 - If vocabulary keeps growing, need to occasionally do the expensive operation of rehashing everything

Trees: Binary Search Tree



Trees: B-tree



Every internal node has a number of children in the interval $[a,b]$ where a, b are appropriate natural numbers, e.g., $[2,4]$.

Trees

- Pros?
 - Solves the prefix problem (terms starting with “ab”)
- Cons?
 - Slower: $O(\log M)$ [and this requires balanced tree]
 - Rebalancing binary trees is expensive
 - But B-trees mitigate the rebalancing problem

Wild-Card Queries: *

- mon*: find all docs containing any word beginning “mon”.
- Easy with binary tree (or B-tree) lexicon
- *mon: find words ending in “mon”: harder
 - Maintain an additional B-tree for terms backwards.
- How can we enumerate all terms meeting the wild-card query pro*cent ?
- Query processing: se*ate AND fil*er ?
 - Expensive

Permuterm Indexes

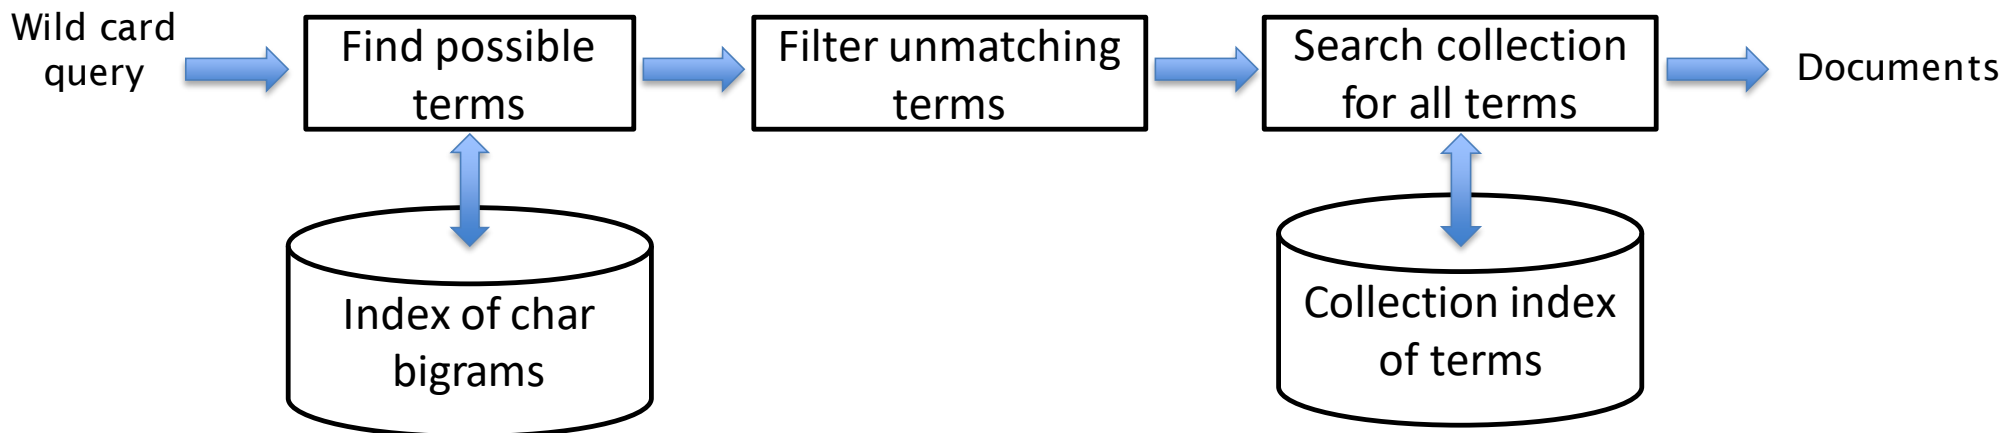

- Transform wild-card queries so that the *s occur at the end
- For term hello, index under:
 - hello\$, ello\$h, llo\$he, lo\$hel, o\$hell, \$hello where \$ is a special symbol.
- Rotate query wild-card to the right
- Queries:
 - X lookup on X\$
 - X* lookup on \$X*
 - *X lookup on X\$*
 - X*Y lookup on Y\$X*
- Index Size?

Character n-gram Indexes

- Enumerate all n-grams (sequence of n chars) occurring in any term
 - e.g., from text “*April is the cruelest month*” we get the 2-grams (bigrams) →
\$a,ap,pr,ri,il,l\$, \$i,is,s\$, \$t,th,he,e\$, \$c,cr,ru,ue,el,le,es,st,t\$,
\$m,mo,on,nt,h\$
 - \$ is a special word boundary symbol
- Maintain a second inverted index from bigrams to dictionary terms that match each bigram.
 - Character n-grams → terms
 - Words → documents

Character n-gram Indexes

- The n -gram index finds *terms* based on a query consisting of n -grams (here $n=2$).



Character n-gram Indexes: Query time

- *Step 1:* Query **mon*** → **\$m AND mo AND on**
 - It would still match **moon**.
- *Step 2:* Must post-filter these terms against query.
 - Phrase match, or post-step1 match
- *Step 3:* Surviving enumerated terms are then looked up in the term-document inverted index.
→ **Montreal OR monster OR monkey**
- Wild-cards can result in expensive query execution (very large disjunctions...)

Character n-gram Indexes: Applications

- Spelling Correction
 - Create n-gram representation for words
 - Build index for words:
 - Dictionary of words → documents (each word is a document)
 - Character n-grams → terms
 - When getting a search term that is misspelled (OOV or not frequent), find possible corrections
 - Possible corrections = most matching results

Query: elepgant → \$e el le ep pg ga an nt t\$

Results:

elegant → \$e el le eg ga an nt t\$

elephant → \$e el le ep ph ha an nt t\$

Character n-gram Indexes: Applications

- Char n-grams can be used as direct index terms for some applications:
 - Arabic IR, when no stemmer/segmenter is available
 - Documents with spelling mistakes: OCR documents
- Word char representation can be with multiple n's
 - “elephant” → 2/3-gram →
“\$e el le ep ph ha an nt t\$ \$el \$ele lep eph pha han ant nt\$”

The **children** behaved well

Her **children** are cute

الأبناء تصرفوا جيدا

أبناءها لطاف

\$ا ال لا أب بن نا اء ع\$

\$أ أب بن نا اء عه ها ا\$

Document: Elepbant → \$e el le ep pb ba an nt t\$

Query: Elephant → \$e el le ep ph ha an nt t\$

Summary

- Index can be multilayer
 - Extent index (multi-terms in one position in document)
- Index does not have to be formed of words
 - Character n-grams representation of words
- Two indexes are sometimes used
 - Index of character n-grams to find matching words
 - Index of terms to search for matched words

Resources

- Text book 1: Intro to IR, Chapter 3.1 – 3.4
- Text book 2: IR in Practice, Chapter 5