



THE UNIVERSITY
of EDINBURGH

Text Technologies for Data Science

INFR11145

Web Search

Instructor:
Walid Magdy

14-Nov-2017


Lecture Objectives

- Learn about:
 - Working with Massive data
 - Link analysis (PageRank)
 - Anchor text



The Web Document Collection

- Huge / Massive
- Graph / Connected
- No design/co-ordination
- Distributed content publishing
- Content includes truth, lies, obsolete information, contradictions ...
- Unstructured (text, html, ...), semi-structured (XML, annotated photos), structured (DB) ...
- Growth – slowed down from initial “volume doubling every few months” but still expanding
- Content can be dynamically generated



Walid Magdy, TTDS 2017/2018



Effect of Massive data


- Web search engines work with huge amount of data
 - 20 PB/day in 2008 → 160 PB/day in 2013 → now??
 - 1 PB = 1,000 TB = 1,000,000 GB
- How this would affect a search engine?
 - Very challenging (storage, processing, networking, ...)
 - Very useful still (makes stuff easier), how?
- Assume two good search engines the collects two sub-sets of the web
 - Search engine A collected N docs → precision@10 = 40%
 - Search engine B collected 4N docs → precision@10??

Walid Magdy, TTDS 2017/2018



Effect of Massive data on Precision

- Assume two good search engines the collects two sub-sets of the web
 - Search engine A collected N docs \rightarrow precision@10 = 40%
 - Search engine B collected $4N$ docs \rightarrow precision@10??
 - Distribution of positive/negative scores stays the same
 - Precision/Recall at a given score stays the same
 - In any decent IR system: more relevant docs exist at the top
 $\rightarrow P@n \uparrow \rightarrow$ precision@10 = 60% (increases)



Big Data or Clever Algorithm?

- For Web search, larger index usually would beat a better retrieval algorithm
 - Google Index vs Bing/Yahoo Index
- Similar to other applications
 - Google MT vs IBM MT
 - Statistical methods trained over **10x** training data beat deep NLP methods with **1x** training data
 - In general ML, the more data, the better the results
 - Tweets classification: using **100x** of noisy training data beats **1x** of well prepared training data, even with absence of stemming & stopping
 - Question answering task:
 - IBM Watson vs Microsoft experiment

Big Data or Clever Algorithm?


- Question answering task:
 - Q: Who created the character of Scrooge?
 - A: Scrooge, introduced by Charles Dickens in "A Christmas Carol"
 - Requires heavy linguistic analysis, lots of research in TREC
- 2002, Microsoft
 - Identify (subj verb obj), rewrite as queries:
 - Q1: "created the character of Scrooge"
 - Q2: "the character of Scrooge was created by"
 - Search the web for exact phrase, get top 500 results
 - Extract phrase: ■ Q1 or Q2 ■ , get most frequent ■
 - Very naive approach, ignores most answers patterns
 - Who cares!! Web is huge, you will find matches anyway

117	Dickens
78	Christmas Carol
75	Charles Dickens
72	Disney
54	Carl Banks
...	

Walid Magdy, TTDS 2017/2018



The Web as a Directed Graph



Assumption 1: A hyperlink between pages denotes author perceived relevance (quality signal)

Assumption 2: The text in the anchor of the hyperlink describes the target page (textual context)

Walid Magdy, TTDS 2017/2018



Links between Pages


- Google Description of **PageRank**:
 - Relies on the “**uniquely democratic**” nature of the web
 - Interprets a link from page A to page B as “**a vote**”
- $A \rightarrow B$: means A thinks B worth something
 - “**wisdom of the crowds**”: many links means B must be good
 - **Content-independent** measure of quality of B
- Use as ranking feature, combined with content
 - But not all pages that link to B are of equal importance!
 - Importance of a link from CNN >>> link from blog page
- Google PageRank, 1998
 - How many “good” pages link to B?

Walid Magdy, TTDS 2017/2018



PageRank: Random Surfer

- Analogy:
 - User starts browsing at a random page
 - Pick a random outgoing link
→ goes there → repeat forever
 - Example:
 $G \rightarrow E \rightarrow F \rightarrow E \rightarrow D \rightarrow B \rightarrow C$
 - With probability $1-\lambda$ jump to a random page
 - Otherwise, can get stuck forever $B \leftrightarrow C$
- **PageRank** of page x
 - Probability of being at page x at a random moment in time




Walid Magdy, TTDS 2017/2018



PageRank: Algorithm

- Initialize $PR_0(x) = \frac{100\%}{N}$
 - N : total number of pages
 - $PR_0(A) = \dots = PR_0(K) = \frac{100\%}{11} = 9.1$



- For every page x

$$PR_{t+1}(x) = \frac{1 - \lambda}{N} + \lambda \sum_{y \rightarrow x} \frac{PR_t(y)}{L_{out}(y)}$$

- $y \rightarrow x$ contributes part of its PR to x
- Spread PR equally among out-links
- Iterate till converge \rightarrow PR scores should sum to 100%


Walid Magdy, TTDS 2017/2018



THE UNIVERSITY
of EDINBURGH

PageRank: Example

- Let $\lambda = 0.82$
- $PR(B) = \frac{0.18}{11} + 0.82 \times [PR(C) + \frac{1}{2}PR(D) + \frac{1}{3}PR(E) + \frac{1}{2}PR(F) + \frac{1}{2}PR(G) + \frac{1}{2}PR(H) + \frac{1}{2}PR(I)]$
 $\approx 0.31 = 31\%$
- $PR(C) = \frac{0.18}{11} + 0.82 \times PR(B)$
 $= 0.18 \times 9.1\% + 0.82 \times 9.1\%$
 $= 9.1\%$
- $PR_{t+1}(C) = 0.18 \times 9.1\% + 0.82 \times 31\%$
 $\approx 26\%$



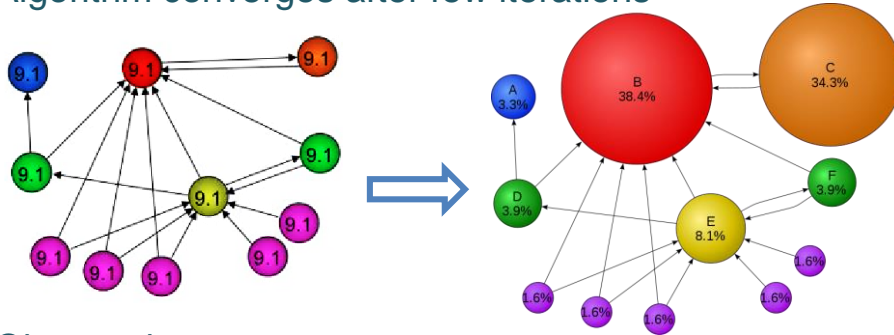
Walid Magdy, TTDS 2017/2018



THE UNIVERSITY
of EDINBURGH

PageRank: Example result

- Algorithm converges after few iterations



- Observations
 - Pages with no inlinks: $PR = (1 - \lambda)/N = 0.18/11 = 1.6\%$
 - Same (or symmetric) inlinks \rightarrow same PR (e.g. **D** and **F**)
 - One inlink from high PR \gg many from low PR (e.g. **C** vs **E**)

Walid Magdy, TTDS 2017/2018



Anchor Text

- Anchor Text (text of a link):
 - Description of destination page
 - Short, descriptive like a query
 - Re-formulated in different ways
 - Human “query expansion”
- Used when indexing page content
 - Add text of all anchor text linking the page
 - Different weights for different anchor text
 - Weighted according to PR of linking page
- Significantly improves retrieval




Walid Magdy, TTDS 2017/2018



Link Spam

- Trackback links (blogs that link to me)
 - Based on `$HTTP_REFERER`
 - Artificial feedback loops
 - Similar to “*follow back*” in Twitter
- Links from comments on sites with high PR
 - Links in comments on CNN
 - One solution: insert `rel=nofollow` into links
 - Link ignored when computing PR
- Link farms
 - Fake densely-connected graph
 - Hundreds of web domains / IPs can be hosted on one machine`



Walid Magdy, TTDS 2017/2018



The Reality

- **PageRank** is used in Google, but is hardly the full story of ranking
 - Many sophisticated features are used
 - Some address specific query classes
 - Machine-learned ranking heavily used
 - Learning to Rank (L2R)
 - Many features are used, including PR
 - Still counted as a very useful feature

Walid Magdy, TTDS 2017/2018



Summary

- Web data is massive
 - Challenging for efficiency, but useful for effectiveness
- PageRank:
 - Probability than random surfer is currently on page x
 - The more powerful pages linking to x, the higher the PR
- Anchor text:
 - Short concise description of target page content
 - Very useful for retrieval
- Link Spam
 - Trackable links, link farms

Resources

- Text book 1: Intro to IR, Section 21.1
- Text Book 2: IR in Practice: 4.5, 10.3
- Page Rank Paper:
Page, L., Brin, S., Motwani, R., & Winograd, T. (1999).
The PageRank citation ranking: Bringing order to the web.
Stanford InfoLab.
- Additional reading:
Dumais, S., Banko, M., Brill, E., Lin, J., & Ng, A. (2002)
Web question answering: Is more always better?.
SIGIR 2002.