

Introductory Applied Machine Learning

Nigel Goddard
School of Informatics

Semester 1

The **primary aim** of the course is to provide the student with a set of practical tools that can be applied to solve real-world problems in machine learning.

Machine learning is the study of computer algorithms that improve automatically through experience [Mitchell, 1997].


In many of today's problems it is

very hard to write a correct program

but very easy to collect examples

Idea behind machine learning:
from the examples, generate the program

Spam Classification



Web page



learning	13
lectures	3
Paris Hilton	0
assignments	7
...	


Feature vector



Classifier

SPAM
NONSPAM

Image Processing




- ▶ Classification: Is there are dog in this image?
- ▶ Localization: If there is a dog in this image, draw its bounding box
- ▶ See: <http://host.robots.ox.ac.uk/pascal/VOC/>

Primate splice-junction gene sequences (DNA)

```
CCAGCTGCATCACAGGAGGCCAGCGAGCAGGTCTGTTCCAAGGGCCTTCGAGCCAGTCTG EI
GAGGTGAAGGACGTCCTTCCCCAGGAGCCGGTGAGAAGCGCAGTCGGGGGCACGGGGATG EI
TAAATTCTTCTGTTTGTTAACACCTTTTCAGACTTATGTGTATGAAGGAGTAGAAGCCAAA IE
AAACTAAAGAATTATTCTTTTACATTTTCAGTTTTTCTTGATCATGAAAACGCCAACAAAA IE
AAAGCAGATCAGCTGTATAAACAGAAAATTATTCGTGGTTTCTGTCACTTGTGTGATGGT N
TTGCCCTCAGCATCACCATGAACGGAGAGGCCATCGCCTGCGCTGAGGGCTGCCAGGCCA N
```

- ▶ Task is to predict if there is an IE (intron/exon), EI or N (neither) junction in the centre of the string
- ▶ Data from ML repository: <http://archive.ics.uci.edu/ml/>

Financial Modeling



News:

Words like **Jackson** and **antitrust** are more likely in the stories preceding the plunge.

$P(\text{shares}) = 0.074$
 $P(\text{antitrust}) = 0.009$
 $P(\text{judge}) = 0.006$
 $P(\text{trading}) = 0.032$
 $P(\text{against}) = 0.025$
 $P(\text{Jackson}) = 0.001$

Software giant Microsoft saw its shares dip a few percentage points this morning after U.S. District Judge Thomas Penfield Jackson issued his "findings of fact" in the government's ongoing antitrust case against the Seattle wealth-creation machine...

$P(\text{shares} \mid \text{MSFT}\downarrow) = 0.071$
 $P(\text{antitrust} \mid \text{MSFT}\downarrow) = 0.044$
 $P(\text{judge} \mid \text{MSFT}\downarrow) = 0.039$
 $P(\text{trading} \mid \text{MSFT}\downarrow) = 0.029$
 $P(\text{against} \mid \text{MSFT}\downarrow) = 0.027$
 $P(\text{Jackson} \mid \text{MSFT}\downarrow) = 0.025$

$$P(\text{MSFT}\downarrow \mid \text{Jackson}) = P(\text{Jackson} \mid \text{MSFT}\downarrow) P(\text{MSFT}\downarrow) / P(\text{Jackson})$$

[Victor Lavrenko]

Collaborative Filtering

Netflix - Unlimited TV Shows & Movies. How Does It Work?

http://www.netflix.com/HowItWorks


Member Sign In

Start Your 1 Month Free Trial How It Works Browse Selection 1 Month Free Trial Info

Unlimited TV episodes & movies instantly over the Internet plus unlimited DVDs by mail!

On your TV

Connect devices like these to your Netflix account to **watch instantly on your TV.**



Wii PS3 XBOX 360

[Learn more >](#) [Learn more >](#) [Learn more >](#)


Watch as often as you want, anytime you want.

[See other devices that stream instantly from Netflix](#)

AND

On your computer


Watch instantly on your computer too!



[Learn more >](#)

PLUS

DVDs by mail



Exchange as often as you want.
No late fees - ever!

FAQs

[How does Netflix work?](#)

[What is the selection like?](#)

[How much does it cost?](#)

[How many DVDs can I rent during my Free Trial?](#)

[How fast will I get my DVDs?](#)

[How long is the free trial?](#)

[Can I cancel any time?](#)

How does Netflix work?

Rent what you want

Simply point and click to add movies & TV episodes to your list. Get DVDs by mail plus instantly watch movies (some new releases) & TV episodes (including current season) online on your PC or Mac or streamed instantly from Netflix over the Internet right to your TV via a Netflix ready device.

New Releases, Classics, TV episodes

Start Your 1 Month Free Trial

Free trial offer details.

Email

Confirm Email

Password 4-10 characters

Confirm Password

More applications

- ▶ Science (Astronomy, neuroscience, medical imaging, bio-informatics)
- ▶ Environment (energy, climate, weather, resources)
- ▶ Retail (Intelligent stock control, demographic store placement)
- ▶ Manufacturing (Intelligent control, automated monitoring, detection methods)
- ▶ Security (Intelligent smoke alarms, fraud detection)
- ▶ Marketing (targetting promotions, ...)
- ▶ Management (Scheduling, timetabling)
- ▶ Finance (credit scoring, risk analysis...)
- ▶ Web data (information retrieval, information extraction, ...)

Overview

- ▶ What is ML? Who uses it?
- ▶ Course structure / Assessment
- ▶ Relationships between ML courses
- ▶ Overview of Machine Learning
- ▶ Overview of the Course
- ▶ Maths Level
- ▶ Reading: W & F chapter 1

Acknowledgements: Thanks to Amos Storkey, David Barber, Chris Williams, Charles Sutton and Victor Lavrenko for permission to use course material from previous years. Additionally, inspiration has been obtained from Geoff Hinton's slides for CSC 2515 in Toronto

Administration

- ▶ All material in course accessible to 3rd- & 4th-year undergraduates. Postgraduates also welcome.
- ▶ Course materials are on Learn (learn.ed.ac.uk)
- ▶ Lectures: online, with quizzes and reviews.
- ▶ Assessment:
 - ▶ Assignments (4) (25% of mark)
 - ▶ Exam (75% of mark)
- ▶ 4(5?) Tutorials and 4+ Labs
- ▶ Course rep
- ▶ Plagiarism
<http://web.inf.ed.ac.uk/infweb/admin/policies/guidelines-plagiarism>
- ▶ Schedule and News

Machine Learning Courses

- IAML** Basic introductory course on supervised and unsupervised learning.
- MLPR** More advanced course on machine learning, including coverage of Bayesian methods.
- RL** Reinforcement Learning.
- MLP** Real-world ML. This year: Deep Learning.
- PMR** Probabilistic modelling and reasoning. Focus on learning and inference for probabilistic models, e.g. probabilistic expert systems, latent variable models, Hidden Markov models
 - Basically, IAML: Users of ML; MLPR: Developers of new ML techniques.

Overview of Machine Learning

- ▶ Supervised learning
 - ▶ Predict an output y when given an input \mathbf{x}
 - ▶ For categorical y : *classification*.
 - ▶ For real-valued y : *regression*.
- ▶ Unsupervised learning
 - ▶ Create an internal representation of the input, e.g. clustering, dimensionality
 - ▶ This is important in machine learning as getting labels is often difficult and expensive
- ▶ Other areas of ML
 - ▶ Learning to predict structured objects (e.g., graphs, trees)
 - ▶ Reinforcement learning (learning from “rewards”)
 - ▶ Semi-supervised learning (combines supervised + unsupervised)
 - ▶ We will not cover these at all in the course

Supervised Learning (Classification)



Feature
processing

Training data

$$\mathbf{x}_1 = (1, 0, 0, 3, \dots)$$

$$y_1 = \text{SPAM}$$

$$\mathbf{x}_2 = (-1, 4, 0, 3, \dots)$$

$$y_2 = \text{NOTSPAM}$$



Learning algorithm

Classifier

Prediction on new
example

$$\mathbf{x}_{1000} = (1, 0, 1, 2, \dots)$$

$$y_{1000} = ???$$

Supervised Learning (Regression)

In this course we will talk about linear regression

$$f(\mathbf{x}) = w_0 + w_1 x_1 + \dots + w_D x_D$$

- ▶ $\mathbf{x} = (x_1, \dots, x_D)^T$
- ▶ Here the assumption is that $f(\mathbf{x})$ is a linear function in \mathbf{x}
- ▶ The specific setting of the parameters w_0, w_1, \dots, w_D is done by minimizing a score function
- ▶ Usual score function is $\sum_{i=1}^n (y^i - f(\mathbf{x}^i))^2$ where the sum runs over all training cases
- ▶ We will cover linear regression later in the course

Unsupervised Learning

In this class we will focus on one kind of unsupervised learning, clustering.



Feature
processing

Training data

$$\mathbf{x}_1 = (1, 0, 0, 3, \dots)$$

$$\mathbf{x}_2 = (-1, 4, 0, 3, \dots)$$

....

$$\mathbf{x}_{1000} = (1, 0, 1, 2, \dots)$$

Cluster labels

$$c_1 = 4$$

$$c_2 = 1$$

....

$$c_2 = 4$$

Learning
algorithm

General structure of supervised learning algorithms

Hand, Mannila, Smyth (2001)

- ▶ Define the **task**
- ▶ Decide on the **model structure** (choice of inductive bias)
- ▶ Decide on the **score function** (judge quality of fitted model)
- ▶ Decide on **optimization/search method** to optimize the score function

Inductive bias

- ▶ Supervised learning is inductive, i.e. we make generalizations about the form of $f(\mathbf{x})$ based on instances \mathcal{D}
- ▶ Let $f(\mathbf{x}; L, \mathcal{D})$ be the function learned by algorithm L with data \mathcal{D}
- ▶ Learning is impossible without making assumptions about f !!

The futility of bias-free learning



1



0



???

The futility of bias-free learning

- ▶ *A learner that makes no a priori assumptions regarding the target concept has no rational basis for classifying any unseen examples* (Mitchell, 1997, p 42)
- ▶ The *inductive bias* of a learner is the set of prior assumptions that it makes (we will not define this formally)
- ▶ We will consider a number of different supervised learning methods in the IAML; these correspond to different inductive biases

Machine Learning and Statistics

- ▶ A lot of work in machine learning can be seen as a rediscovery of things that were known in statistics; but there are also flows in the other direction
- ▶ The emphasis is rather different. One difference is a focus on *prediction* in machine learning vs *interpretation* of the model in statistics
- ▶ Until recently, machine learning usually referred to tasks associated with artificial intelligence (AI) such as recognition, diagnosis, planning, robot control, prediction, etc. These provide rich and interesting tasks
- ▶ Today interesting machine learning tasks abound.
- ▶ Goals can be autonomous machine performance, or enabling humans to learn from data (data mining).

Provisional Course Outline

- ▶ Introduction (today)
- ▶ Basic probability
- ▶ Thinking about data
- ▶ Naïve Bayes classification
- ▶ Decision trees
- ▶ Linear regression
- ▶ Generalization and Overfitting
- ▶ Linear classification: logistic regression, perceptrons
- ▶ Kernel classifiers: support vector machines
- ▶ Dimensionality reduction (PCA etc)
- ▶ Performance evaluation
- ▶ Clustering (k -means, hierarchical)
- ▶ Neural Networks

Maths Level

- ▶ Machine learning generally involves a significant number of mathematical ideas and a significant amount of mathematical manipulation
- ▶ IAML aims to keep the maths level to a minimum, explaining things more in terms of higher-level concepts, and developing understanding in a procedural way (e.g. how to program an algorithm)
- ▶ For those wanting to pursue research in any of the areas covered you will need courses like PMR, MLPR & MLP