

# IAML: Linear Regression

Nigel Goddard  
School of Informatics

Semester 1

- ▶ The linear model
- ▶ Fitting the linear model to data
- ▶ Probabilistic interpretation of the error function
- ▶ Examples of regression problems
- ▶ Dealing with multiple outputs
- ▶ Generalized linear regression
- ▶ Radial basis function (RBF) models

# The Regression Problem

- ▶ Classification and regression problems:
  - ▶ Classification: target of prediction is discrete
  - ▶ Regression: target of prediction is continuous
- ▶ Training data: Set  $\mathcal{D}$  of pairs  $(\mathbf{x}_i, y_i)$  for  $i = 1, \dots, n$ , where  $\mathbf{x}_i \in \mathbb{R}^D$  and  $y_i \in \mathbb{R}$
- ▶ Today: Linear regression, i.e., relationship between  $\mathbf{x}$  and  $y$  is linear.
- ▶ Although this is simple (and limited) it is:
  - ▶ More powerful than you would expect
  - ▶ The basis for more complex nonlinear methods
  - ▶ Teaches a lot about regression and classification

# Examples of regression problems

- ▶ Robot inverse dynamics: predicting what torques are needed to drive a robot arm along a given trajectory
- ▶ Electricity load forecasting, generate hourly forecasts two days in advance.
- ▶ Predicting staffing requirements at help desks based on historical data and product and sales information,
- ▶ Predicting the time to failure of equipment based on utilization and environmental conditions

# The Linear Model

- ▶ Linear model


$$\begin{aligned}f(\mathbf{x}; \mathbf{w}) &= w_0 + w_1 x_1 + \dots + w_D x_D \\ &= \phi(\mathbf{x})\mathbf{w}\end{aligned}$$

where  $\phi(\mathbf{x}) = (1, x_1, \dots, x_D) = (1, \mathbf{x}^T)$   
and



$$\mathbf{w} = \begin{pmatrix} w_0 \\ w_1 \\ \dots \\ w_D \end{pmatrix} \tag{1}$$

- ▶ The maths of fitting linear models to data is easy. We use the notation  $\phi(\mathbf{x})$  to make generalisation easy later.


# Toy example: Data



# Toy example: Data



# With two features



Instead of a line, a *plane*. With more features, a *hyperplane*.



## CPU Performance Data Set

- ▶ Predict: PRP: published relative performance
- ▶ MYCT: machine cycle time in nanoseconds (integer)
- ▶ MMIN: minimum main memory in kilobytes (integer)
- ▶ MMAX: maximum main memory in kilobytes (integer)
- ▶ CACH: cache memory in kilobytes (integer)
- ▶ CHMIN: minimum channels in units (integer)
- ▶ CHMAX: maximum channels in units (integer)

## With more features

```
PRP = - 56.1  
      + 0.049 MYCT  
      + 0.015 MMIN  
      + 0.006 MMAX  
      + 0.630 CACH  
      - 0.270 CHMIN  
      + 1.46  CHMAX
```

# In matrix notation

- ▶ Design matrix is  $n \times (D + 1)$

$$\Phi = \begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1D} \\ 1 & x_{21} & x_{22} & \dots & x_{2D} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} & \dots & x_{nD} \end{pmatrix}$$

- ▶  $x_{ij}$  is the  $j$ th component of the training input  $\mathbf{x}_i$
- ▶ Let  $\mathbf{y} = (y_1, \dots, y_n)^T$
- ▶ Then  $\hat{\mathbf{y}} = \Phi \mathbf{w}$  is ...?

# Linear Algebra: The 1-Slide Version

What is matrix multiplication?

$$A = \begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{pmatrix}, \mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \\ b_3 \end{pmatrix}$$

First consider matrix times vector, i.e.,  $\mathbf{Ab}$ . Two answers:

1.  $\mathbf{Ab}$  is a linear combination of the columns of  $A$

$$\mathbf{Ab} = b_1 \begin{pmatrix} a_{11} \\ a_{21} \\ a_{31} \end{pmatrix} + b_2 \begin{pmatrix} a_{12} \\ a_{22} \\ a_{32} \end{pmatrix} + b_3 \begin{pmatrix} a_{13} \\ a_{23} \\ a_{33} \end{pmatrix}$$

2.  $\mathbf{Ab}$  is a vector. Each element of the vector is the dot products between  $b$  and *one row* of  $A$ .

$$\mathbf{Ab} = \begin{pmatrix} (a_{11}, a_{12}, a_{13})\mathbf{b} \\ (a_{21}, a_{22}, a_{23})\mathbf{b} \\ (a_{31}, a_{32}, a_{33})\mathbf{b} \end{pmatrix}$$

## Linear model (part 2)

In matrix notation:

- ▶ Design matrix is  $n \times (D + 1)$

$$\Phi = \begin{pmatrix} 1 & x_{11} & x_{12} & \dots & x_{1D} \\ 1 & x_{21} & x_{22} & \dots & x_{2D} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} & \dots & x_{nD} \end{pmatrix}$$

- ▶  $x_{ij}$  is the  $j$ th component of the training input  $\mathbf{x}_i$
- ▶ Let  $\mathbf{y} = (y_1, \dots, y_n)^T$
- ▶ Then  $\hat{\mathbf{y}} = \Phi \mathbf{w}$  is the model's predicted values on training inputs.

# Solving for Model Parameters

This looks like what we've seen in linear algebra

$$\mathbf{y} = \Phi \mathbf{w}$$

We know  $\mathbf{y}$  and  $\Phi$  but not  $\mathbf{w}$ .

So why not take  $\mathbf{w} = \Phi^{-1} \mathbf{y}$ ? (You can't, but why?)

# Solving for Model Parameters

This looks like what we've seen in linear algebra

$$\mathbf{y} = \Phi \mathbf{w}$$

We know  $\mathbf{y}$  and  $\Phi$  but not  $\mathbf{w}$ .

So why not take  $\mathbf{w} = \Phi^{-1} \mathbf{y}$ ? (You can't, but why?)

Three reasons:

- ▶  $\Phi$  is not square. It is  $n \times (D + 1)$ .
- ▶ The system is overconstrained ( $n$  equations for  $D + 1$  parameters), in other words
- ▶ The data has noise


Want a *loss function*  $O(\mathbf{w})$  that

- ▶ We minimize wrt  $\mathbf{w}$ .
- ▶ At minimum,  $\hat{\mathbf{y}}$  looks like  $\mathbf{y}$ .
- ▶ (Recall:  $\hat{\mathbf{y}}$  depends on  $\mathbf{w}$ )

$$\hat{\mathbf{y}} = \Phi \mathbf{w}$$



# Fitting a linear model to data



- ▶ A common choice: *squared error* (makes the maths easy)

$$O(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$


- ▶ In the picture: this is sum of squared length of black sticks.
- ▶ (Each one is called a *residual*, i.e., each  $y_i - \mathbf{w}^T \mathbf{x}_i$ )

# Fitting a linear model to data



$$\begin{aligned}O(\mathbf{w}) &= \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2 \\ &= (\mathbf{y} - \Phi \mathbf{w})^T (\mathbf{y} - \Phi \mathbf{w})\end{aligned}$$

- ▶ We want to minimize this with respect to  $\mathbf{w}$ .
- ▶ The error surface is a parabolic bowl



- ▶ How do we do this?

# The Solution

- ▶ Answer: to minimize  $O(\mathbf{w}) = \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$ , set partial derivatives to 0.
- ▶ This has an analytical solution

$$\hat{\mathbf{w}} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{y}$$

- ▶  $(\Phi^T \Phi)^{-1} \Phi^T$  is the pseudo-inverse of  $\Phi$
- ▶ First check: Does this make sense? Do the matrix dimensions line up?
- ▶ Then: Why is this called a pseudo-inverse? ()
- ▶ Finally: What happens if there are no features?

# Probabilistic interpretation of $O(\mathbf{w})$

- ▶ Assume that  $y = \mathbf{w}^T \mathbf{x} + \epsilon$ , where  $\epsilon \sim N(0, \sigma^2)$
- ▶ (This is an exact linear relationship plus Gaussian noise.)
- ▶ This implies that  $y|\mathbf{x}_i \sim N(\mathbf{w}^T \mathbf{x}_i, \sigma^2)$ , i.e.

$$-\log p(y_i|\mathbf{x}_i) = \log \sqrt{2\pi} + \log \sigma + \frac{(y_i - \mathbf{w}^T \mathbf{x}_i)^2}{2\sigma^2}$$

- ▶ So minimising  $O(\mathbf{w})$  equivalent to maximising likelihood!
- ▶ Can view  $\mathbf{w}^T \mathbf{x}$  as  $E[y|\mathbf{x}]$ .
- ▶ Squared residuals allow estimation of  $\sigma^2$


$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

Fitting this into the general structure for learning algorithms:

- ▶ Define the **task**: regression
- ▶ Decide on the **model structure**: linear regression model
- ▶ Decide on the **score function**: squared error (likelihood)
- ▶ Decide on **optimization/search method** to optimize the score function: calculus (analytic solution)

# Sensitivity to Outliers

- ▶ Linear regression is sensitive to *outliers*
- ▶ Example: Suppose  $y = 0.5x + \epsilon$ , where  $\epsilon \sim N(0, \sqrt{0.25})$ , and then add a point (2.5,3):



Graphical diagnostics can be useful for checking:

- ▶ Is the relationship obviously nonlinear? Look for structure in residuals?
- ▶ Are there obvious outliers?

The goal isn't to find all problems. You can't. The goal is to find obvious, embarrassing problems.

Examples: Plot residuals by fitted values. Stats packages will do this for you.


# Dealing with multiple outputs

- ▶ Suppose there are  $q$  different targets for each input  $\mathbf{x}$
- ▶ We introduce a different  $\mathbf{w}_i$  for each target dimension, and do regression separately for each one
- ▶ This is called *multiple regression*




# Basis expansion


- We can easily transform the original attributes  $\mathbf{x}$  non-linearly into  $\phi(\mathbf{x})$  and do linear regression on them



polynomial



Gaussians



sigmoids

Figure credit: Chris Bishop, PRML

- ▶ Design matrix is  $n \times m$

$$\Phi = \begin{pmatrix} \phi_1(\mathbf{x}_1) & \phi_2(\mathbf{x}_1) & \dots & \phi_m(\mathbf{x}_1) \\ \phi_1(\mathbf{x}_2) & \phi_2(\mathbf{x}_2) & \dots & \phi_m(\mathbf{x}_2) \\ \vdots & \vdots & \vdots & \vdots \\ \phi_1(\mathbf{x}_n) & \phi_2(\mathbf{x}_n) & \dots & \phi_m(\mathbf{x}_n) \end{pmatrix}$$





- ▶ Let  $\mathbf{y} = (y_1, \dots, y_n)^T$
- ▶ Minimize  $E(\mathbf{w}) = \|\mathbf{y} - \Phi\mathbf{w}\|^2$ . As before we have an analytical solution

$$\hat{\mathbf{w}} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{y}$$

- ▶  $(\Phi^T \Phi)^{-1} \Phi^T$  is the pseudo-inverse of  $\Phi$

# Example: polynomial regression

$$\phi(x) = (1, x, x^2, \dots, x^M)^T$$



# More about the features

- ▶ Transforming the features can be important.
- ▶ Example: Suppose I want to predict the CPU performance.
- ▶ Maybe one of the features is *manufacturer*.

$$x_1 = \begin{cases} 1 & \text{if Intel} \\ 2 & \text{if AMD} \\ 3 & \text{if Apple} \\ 4 & \text{if Motorola} \end{cases}$$

- ▶ Let's use this as a feature. Will this work?

# More about the features

- ▶ Transforming the features can be important.
- ▶ Example: Suppose I want to predict the CPU performance.
- ▶ Maybe one of the features is *manufacturer*.

$$x_1 = \begin{cases} 1 & \text{if Intel} \\ 2 & \text{if AMD} \\ 3 & \text{if Apple} \\ 4 & \text{if Motorola} \end{cases}$$

- ▶ Let's use this as a feature. Will this work?
- ▶ Not the way you want. Do you really believe AMD is double Intel?

- ▶ Instead convert this into 0/1

$x_1 = 1$  if Intel, 0 otherwise

$x_2 = 1$  if AMD, 0 otherwise


$\vdots$

- ▶ Note this is a consequence of linearity. We saw something similar with text in the first week.
- ▶ Other good transformations: log, square, square root

# Radial basis function (RBF) models



- ▶ Set  $\phi_i(\mathbf{x}) = \exp(-\frac{1}{2}|\mathbf{x} - \mathbf{c}_i|^2/\alpha^2)$ .
- ▶ Need to position these “basis functions” at some prior chosen centres  $\mathbf{c}_i$  and with a given width  $\alpha$ . There are many ways to do this but choosing a subset of the datapoints as centres is one method that is quite effective
- ▶ Finding the weights is the same as ever: the pseudo-inverse solution.

# RBF example






# RBF example




# An RBF feature

Original data




RBF feature,  $c_1 = 3$ ,  $\alpha_1 = 1$




# Another RBF feature

Notice how the feature functions “specialize” in input space.

Original data



RBF feature,  $c_2 = 6$ ,  $\alpha_2 = 1$




# RBF example


Run the RBF with both basis functions above, plot the residuals

$$y_i - \phi(\mathbf{x}_i)^T \mathbf{w}$$

Original data



Residuals



# RBF: Ay, there's the rub

- ▶ So why not use RBFs for everything?
- ▶ Short answer: You might need too many basis functions.
- ▶ This is especially true in high dimensions (we'll say more later)
- ▶ Too many means you probably overfit.
- ▶ Extreme example: Centre one on each training point.
- ▶ Also: notice that we haven't seen yet in the course how to learn the RBF *parameters*, i.e., the mean and standard deviation of each kernel
- ▶ Main point of presenting RBFs now: Set up later methods like support vector machines

# Summary

- ▶ Linear regression often useful out of the box.
- ▶ More useful than it would be seem because linear means linear *in the parameters*. You can do a nonlinear transform of the data first, e.g., polynomial, RBF. This point will come up again.
- ▶ Maximum likelihood solution is computationally efficient (pseudo-inverse)
- ▶ Danger of overfitting, especially with many features or basis functions