

# IAML: Dimensionality Reduction

Victor Lavrenko and Nigel Goddard  
School of Informatics


Semester 1

# Overview

- Curse of dimensionality
- Different ways to reduce dimensionality
- Principal Components Analysis (PCA)
- Example: Eigen Faces
- PCA for classification
- Witten & Frank section 7.3
  - only the PCA section required

# True vs. observed dimensionality

- Get a population, predict some property
  - instances represented as {urefu, height} pairs
  - what is the dimensionality of this data?




- Data points over time from different geographic areas over time:
  - $X_1$ : # of skidding accidents
  - $X_2$ : # of burst water pipes
  - $X_3$ : snow-plow expenditures
  - $X_4$ : # of school closures
  - $X_5$ : # patients with heat stroke

Temperature?

# Curse of dimensionality

- Datasets typically high dimensional
  - vision:  $10^4$  pixels, text:  $10^6$  words
    - the way we observe / record them
  - true dimensionality often much lower
    - a manifold (sheet) in a high-d space
- Example: handwritten digits
  - 28 x 28 bitmap:  $\{0,1\}^{784}$  possible events
    - will never see most of these events
    - actual digits: tiny fraction of events
  - true dimensionality:
    - possible variations of the pen-stroke



# Curse of dimensionality (2)

- Machine learning methods are statistical by nature
  - count observations in various regions of some space
  - use counts to construct the predictor  $f(x)$
  - e.g. decision trees:  $p_+/p_-$  in  $\{o=\text{rain}, w=\text{strong}, T>28^\circ\}$
  - text: #documents in  $\{\text{"hp" and "3d" and not "S" and ...}\}$
- As dimensionality grows: fewer observations per region
  - 1d: 3 regions, 2d:  $3^2$  regions, 1000d – hopeless
  - statistics need repetition
    - flip a coin once  $\rightarrow$  head
    - $P(\text{head}) = 100\%$ ?




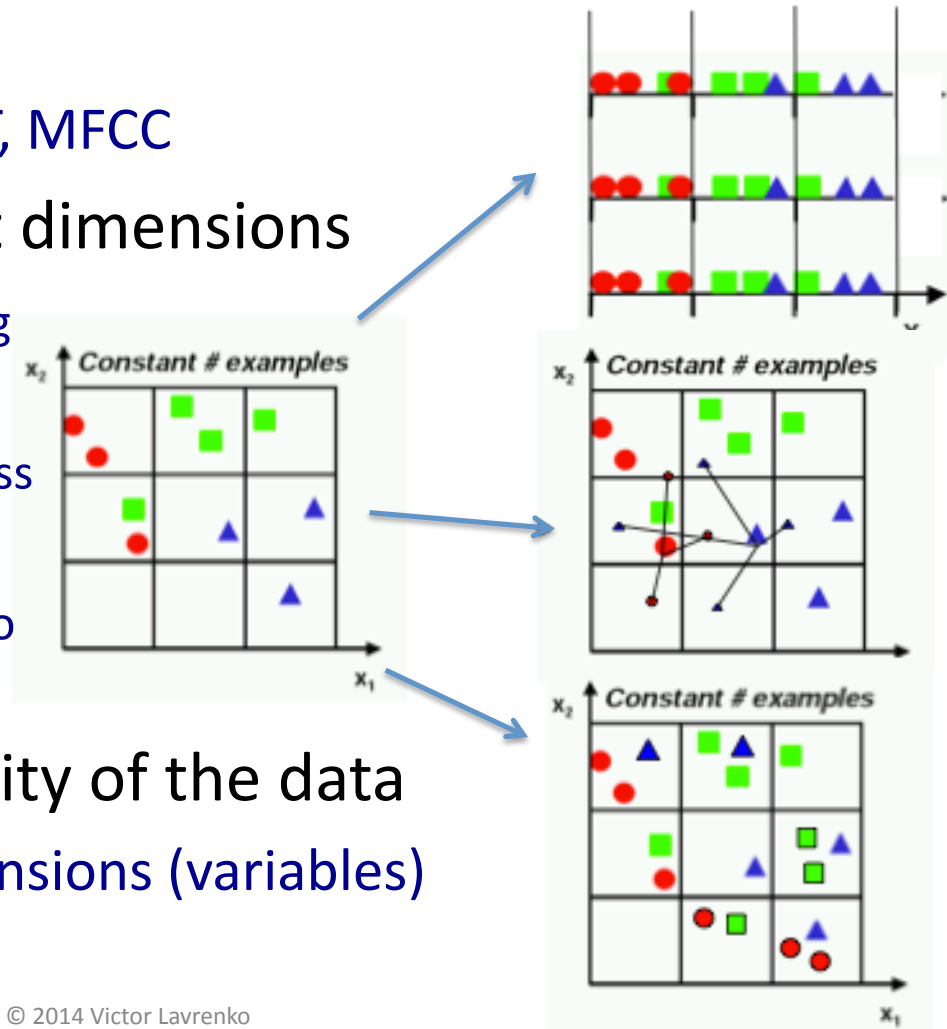


Figure © Ricardo Gutierrez-Osuna


# Dealing with high dimensionality

- Use domain knowledge
  - feature engineering: SIFT, MFCC
- Make assumption about dimensions
  - independence: count along each dimension separately
  - smoothness: propagate class counts to neighboring regions
  - symmetry: e.g. invariance to order of dimensions:  $x_1 \Leftrightarrow x_2$
- Reduce the dimensionality of the data
  - create a new set of dimensions (variables)






# Dimensionality reduction

- Goal: represent instances with fewer variables
  - try to preserve as much structure in the data as possible
  - discriminative: only structure that affects class separability
- Feature selection
  - pick a subset of the original dimensions  $X_1$   $X_2$   $X_3 \dots X_{d-1}$   $X_d$
  - discriminative: pick good class “predictors” (e.g. gain)
- Feature extraction
  - construct a new set of dimensions  $E_1$   $E_2 \dots E_m$   
 $E_i = f(X_1 \dots X_d)$
  - (linear) combinations of original  $X_1$   $X_2$   $X_3 \dots X_d$



# Principal Components Analysis

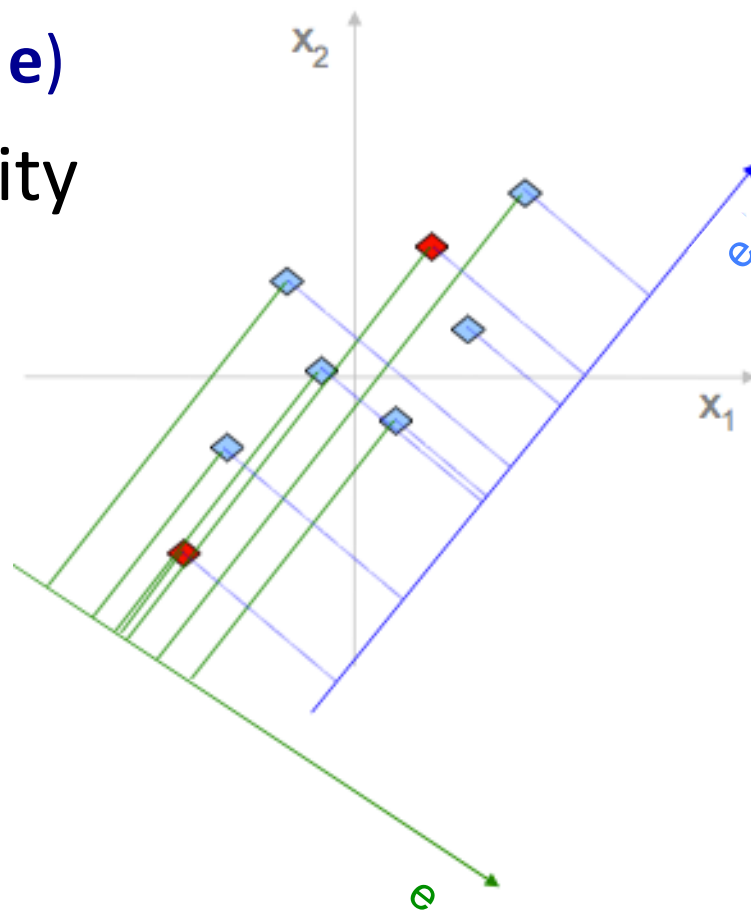
- Defines a set of principal components
  - 1<sup>st</sup>: direction of the greatest variability in the data
  - 2<sup>nd</sup>: perpendicular to 1<sup>st</sup>, greatest variability of what's left
  - ... and so on until  $d$  (original dimensionality)
- First  $m \ll d$  components become  $m$  new dimensions
  - change coordinates of every data point to these dimensions






# Why greatest variability?

- Example: reduce 2-dimensional data to 1-d
  - $\{x_1, x_2\} \rightarrow e'$  (along new axis  $e$ )
- Pick  $e$  to maximize variability
- Reduces cases when two points are close in  $e$ -space but very far in  $(x, y)$ -space
- Minimizes distances between original points and their projections



# Principal components

- “Center” the data at zero:  $x_{i,a} = x_{i,a} - \mu_a$ 
  - subtract mean from each attribute
- Compute covariance matrix  $\Sigma$ 
  - covariance of dimensions  $x_1$  and  $x_2$ :
    - do  $x_1$  and  $x_2$  tend to increase together?
    - or does  $x_2$  decrease as  $x_1$  increases?
- Multiply a vector by  $\Sigma$ :  $\begin{pmatrix} 2.0 & 0.8 \\ 0.8 & 0.6 \end{pmatrix} \begin{pmatrix} -1 \\ +1 \end{pmatrix} \rightarrow \begin{pmatrix} -1.2 \\ -0.2 \end{pmatrix}$  again  $\rightarrow \begin{pmatrix} -2.5 \\ -1.0 \end{pmatrix} \rightarrow \begin{pmatrix} -6.0 \\ -2.7 \end{pmatrix} \rightarrow \begin{pmatrix} -14.1 \\ -6.4 \end{pmatrix} \rightarrow \begin{pmatrix} -33.3 \\ -15.1 \end{pmatrix}$ 
  - turns towards direction of variance
- Want vectors  $\mathbf{e}$  which aren't turned:  $\Sigma \mathbf{e} = \lambda \mathbf{e}$ 
  - $\mathbf{e}$  ... eigenvectors of  $\Sigma$ ,  $\lambda$  ... corresponding eigenvalues
  - principal components = eigenvectors w. largest eigenvalues



# Finding Principal Components

1. find eigenvalues by solving:  $\det(\Sigma - \lambda I) = 0$

$$\det \begin{pmatrix} 2.0 - \lambda & 0.8 \\ 0.8 & 0.6 - \lambda \end{pmatrix} = (2 - \lambda)(0.6 - \lambda) - (0.8)(0.8) = \lambda^2 - 2.6\lambda + 0.56 = 0$$
$$\{\lambda_1, \lambda_2\} = \frac{1}{2} \left( 2.6 \pm \sqrt{2.6^2 - 4 * 0.56} \right) = \{2.36, 0.23\}$$

2. find  $i^{\text{th}}$  eigenvector by solving:  $\Sigma \mathbf{e}_i = \lambda_i \mathbf{e}_i$

$$\begin{pmatrix} 2.0 & 0.8 \\ 0.8 & 0.6 \end{pmatrix} \begin{pmatrix} e_{1,1} \\ e_{1,2} \end{pmatrix} = 2.36 \begin{pmatrix} e_{1,1} \\ e_{1,2} \end{pmatrix} \Rightarrow \begin{cases} 2.0e_{1,1} + 0.8e_{1,2} = 2.36e_{1,1} \\ 0.8e_{1,1} + 0.6e_{1,2} = 2.36e_{1,2} \end{cases} \Rightarrow e_{1,1} = 2.2e_{1,2}$$

$$\begin{pmatrix} 2.0 & 0.8 \\ 0.8 & 0.6 \end{pmatrix} \begin{pmatrix} e_{2,1} \\ e_{2,2} \end{pmatrix} = 0.23 \begin{pmatrix} e_{2,1} \\ e_{2,2} \end{pmatrix} \Rightarrow e_2 = \begin{bmatrix} -0.41 \\ 0.91 \end{bmatrix}$$

$$e_1 \sim \begin{bmatrix} 2.2 \\ 1 \end{bmatrix}$$


want:  $\|e_1\| = 1$

$$e_1 = \begin{bmatrix} 0.91 \\ 0.41 \end{bmatrix}$$

3. 1<sup>st</sup> PC:  $\begin{bmatrix} 0.91 \\ 0.41 \end{bmatrix}$ , 2<sup>nd</sup> PC:  $\begin{bmatrix} -0.41 \\ 0.91 \end{bmatrix}$

# Projecting to new dimensions

- $\mathbf{e}_1 \dots \mathbf{e}_m$  are new dimension vectors
- Have instance  $\mathbf{x} = \{x_1 \dots x_d\}$  (original coordinates)
- Want new coordinates  $\mathbf{x}' = \{x'_1 \dots x'_m\}$ :
  1. “center” the instance (subtract the mean):  $\mathbf{x}' - \boldsymbol{\mu}$
  2. “project” to each dimension:  $(\mathbf{x}' - \boldsymbol{\mu})^T \mathbf{e}_j$  for  $j=1 \dots m$




$$(\vec{x} - \vec{\mu}) = \begin{bmatrix} (x_1 - \mu_1) & (x_2 - \mu_2) & \cdots & (x_d - \mu_d) \end{bmatrix}$$
$$\begin{bmatrix} x'_1 \\ x'_2 \\ \vdots \\ x'_m \end{bmatrix} = \begin{bmatrix} (\vec{x} - \vec{\mu})^T \vec{e}_1 \\ (\vec{x} - \vec{\mu})^T \vec{e}_2 \\ \vdots \\ (\vec{x} - \vec{\mu})^T \vec{e}_m \end{bmatrix} = \begin{bmatrix} (x_1 - \mu_1)e_{1,1} + (x_2 - \mu_2)e_{1,2} + \cdots + (x_d - \mu_d)e_{1,d} \\ (x_1 - \mu_1)e_{2,1} + (x_2 - \mu_2)e_{2,2} + \cdots + (x_d - \mu_d)e_{2,d} \\ \vdots \\ (x_1 - \mu_1)e_{m,1} + (x_2 - \mu_2)e_{m,2} + \cdots + (x_d - \mu_d)e_{m,d} \end{bmatrix}$$

# Direction of greatest variability

- Select dimension  $\mathbf{e}$  which maximizes the variance

- Points  $\mathbf{x}_i$  “projected” onto vector  $\mathbf{e}$ :

- Variance of projections:  $\frac{1}{n} \sum_{i=1}^n \left( \sum_{j=1}^d x_{ij} e_j - \mu \right)^2 = \frac{1}{n} \sum_{i=1}^n \left( \sum_{j=1}^d x_{ij} e_j \right)^2$



- Maximize variance
  - want unit length:  $\|\mathbf{e}\|=1$
  - add Lagrange multiplier

$$V = \frac{1}{n} \sum_{i=1}^n \left( \sum_{j=1}^d x_{ij} e_j \right)^2 - \lambda \left( \left( \sum_{k=1}^d e_k^2 \right) - 1 \right)$$

$$\frac{\partial V}{\partial e_a} = \frac{2}{n} \sum_{i=1}^n \left( \sum_{j=1}^d x_{ij} e_j \right) x_{ia} - 2\lambda e_a = 0$$

$$\Sigma \mathbf{e} = \lambda \mathbf{e} \quad \left\{ \begin{array}{l} \sum_{j=1}^d \text{cov}(1,j) e_j = \lambda e_1 \\ \vdots \\ \sum_{j=1}^d \text{cov}(d,j) e_j = \lambda e_d \end{array} \right.$$


$\mathbf{e}$  must be an eigenvector

hold for  
 $a=1..d$

$$2 \sum_{j=1}^d e_j \underbrace{\left( \frac{1}{n} \sum_{i=1}^n x_{ia} x_{ij} \right)}_{\text{covariance of } a,j} = 2\lambda e_a$$

# Variance along eigenvector

Variance of projected points ( $\mathbf{x}^T \mathbf{e}$ ):



$$\begin{aligned}
 \frac{1}{n} \sum_{i=1}^n \left( \sum_{j=1}^d x_{ij} \mathbf{e}_j - \mu \right)^2 &= \frac{1}{n} \sum_{i=1}^n \left( \sum_{j=1}^d x_{ij} \mathbf{e}_j \right)^2 \\
 &= \frac{1}{n} \sum_{i=1}^n \left( \sum_{j=1}^d x_{ij} \mathbf{e}_j \right) \left( \sum_{a=1}^d x_{ia} \mathbf{e}_a \right) \\
 &= \sum_{a=1}^d \sum_{j=1}^d \left( \frac{1}{n} \sum_{i=1}^n x_{ia} x_{ij} \right) \mathbf{e}_j \mathbf{e}_a \\
 &= \sum_{a=1}^d \left( \sum_{j=1}^d \text{cov}(a, j) \mathbf{e}_j \right) \mathbf{e}_a \\
 &= \sum_{a=1}^d (\lambda \mathbf{e}_a) \mathbf{e}_a \\
 &= \lambda \|\mathbf{e}\|^2 = \lambda
 \end{aligned}$$



$$\begin{aligned}
 \mu &= \frac{1}{n} \sum_{i=1}^n \left( \sum_{j=1}^d x_{ij} \mathbf{e}_j \right) \\
 &= \sum_{j=1}^d \left( \frac{1}{n} \sum_{i=1}^n x_{ij} \right) \mathbf{e}_j
 \end{aligned}$$




$$\text{cov}(a, j) = \frac{1}{n} \sum_{i=1}^n x_{ia} x_{ij}$$



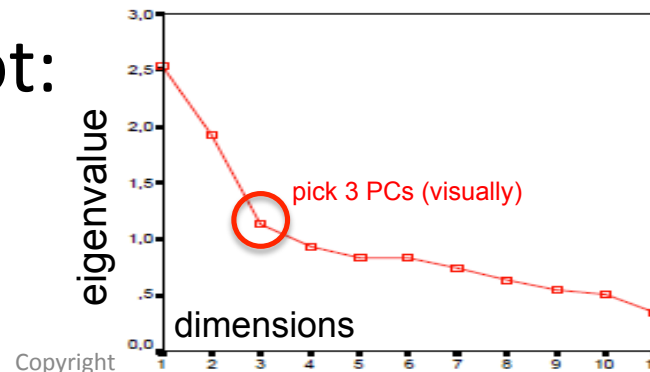
$$\sum_{j=1}^d \text{cov}(a, j) \mathbf{e}_j = \lambda \mathbf{e}_a \quad \text{\textcolor{blue}{e is an eigenvector of the covariance matrix}}$$

# How many dimensions?

- Have: eigenvectors  $\mathbf{e}_1 \dots \mathbf{e}_d$  want:  $m \ll d$
- Proved: eigenvalue  $\lambda_i$  = variance along  $\mathbf{e}_i$
- Pick  $\mathbf{e}_i$  that “explain” the most variance
  - sort eigenvectors s.t.  $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d$
  - pick first  $m$  eigenvectors which explain 90% or the total variance
    - typical threshold values: 0.9 or 0.95




- Or use a scree plot:
  - like K-means




# PCA in a nutshell

## 1. correlated hi-d data

("urefu" means "height" in Swahili)



## 2. center the points



## 3. compute covariance matrix

$$\begin{matrix} & h & u \\ h & \begin{pmatrix} 2.0 & 0.8 \end{pmatrix} \\ u & \begin{pmatrix} 0.8 & 0.6 \end{pmatrix} \end{matrix} \rightarrow \text{cov}(h,u) = \frac{1}{n} \sum_{i=1}^n h_i u_i$$


## 4. eigenvectors + eigenvalues

$$\begin{pmatrix} 2.0 & 0.8 \\ 0.8 & 0.6 \end{pmatrix} \begin{pmatrix} e_h \\ e_u \end{pmatrix} = \lambda_e \begin{pmatrix} e_h \\ e_u \end{pmatrix}$$

$$\begin{pmatrix} 2.0 & 0.8 \\ 0.8 & 0.6 \end{pmatrix} \begin{pmatrix} f_h \\ f_u \end{pmatrix} = \lambda_f \begin{pmatrix} f_h \\ f_u \end{pmatrix}$$

`eig(cov(data))`

## 5. pick m < d eigenvectors w. highest eigenvalues




## 6. project data points to those eigenvectors

$$x'_e = x^T e = \sum_{j=1}^d x_{ij} e_j$$

A diagram showing a data point (red dot) being projected onto the principal axis 'e' (green line). The projection is shown as a red dot on the axis. The axes are labeled 'e' and 'f'. The data points are shown as green diamonds.


## 7. uncorrelated low-d data





# PCA example: Eigen Faces

input: dataset of  $N$  face images




face:  $K \times K$  bitmap of pixels



“unfold” each bitmap to  $K^2$ -dimensional vector

arrange in a matrix  
each face = column




Matlab demo on course webpage

can visualize  
eigenvectors:  
 $m$  “aspects”  
of prototypical  
facial features




“fold” into a  $K \times K$  bitmap



PCA

$K^2 \times m$

set of  $m$  eigenvectors  
each is  $K^2$ -dimensional




# Eigen Faces: Projection

The equation is: 
$$\text{Target Face} = \text{mean} + 0.9 * \text{Eigenface 1} - 0.2 * \text{Eigenface 2} + 0.4 * \text{Eigenface 3} + \dots$$



- Project new face to space of eigen-faces
- Represent vector as a linear combination of principal components
- How many do we need?



# (Eigen) Face Recognition

- Face similarity
  - in the reduced space
  - insensitive to lighting expression, orientation
- Projecting new “faces”
  - everything is a face




new face

projected to eigenfaces


# PCA: practical issues

- Covariance extremely sensitive to large values
  - multiply some dimension by 1000
    - dominates covariance
    - becomes a principal component
  - normalize each dimension to zero mean and unit variance:  
 $\mathbf{x}' = (\mathbf{x} - \text{mean}) / \text{st.dev}$
- PCA assumes underlying subspace is linear
  - 1d: straight line
  - 2d: flat sheet
  - transform to handle non-linear spaces (manifolds)




# PCA and classification

- PCA is unsupervised
  - maximizes overall variance of the data along a small set of directions
  - does not know anything about class labels
  - can pick direction that makes it hard to separate classes
- Discriminative approach
  - look for a dimension that makes it easy to separate classes




# Linear Discriminant Analysis

- LDA: pick a new dimension that gives:
  - maximum separation between means of projected classes
  - minimum variance within each projected class
- Solution: eigenvectors based on between-class and within-class covariance matrices



# PCA vs. LDA

- LDA not guaranteed to be better for classification
  - assumes classes are unimodal Gaussians
  - fails when discriminatory information is not in the mean, but in the variance of the data
- Example where PCA gives a better projection:



# Dimensionality reduction

- Pros
  - reflects our intuitions about the data
  - allows estimating probabilities in high-dimensional data
    - no need to assume independence etc.
  - dramatic reduction in size of data
    - faster processing (as long as reduction is fast), smaller storage
- Cons
  - too expensive for many applications (Twitter, web)
  - disastrous for tasks with fine-grained classes
  - understand assumptions behind the methods (linearity etc.)
    - there may be better ways to deal with sparseness



# Summary

- True dimensionality  $\ll$  observed dimensionality
- High dimensionality  $\rightarrow$  sparse, unstable estimates
- Dealing with high dimensionality:
  - use domain knowledge
  - make an assumption: independence / smoothness / symmetry
  - dimensionality reduction: feature selection / feature extraction
- Principal Components Analysis (PCA)
  - picks dimensions that maximize variability
    - eigenvectors of the covariance matrix
  - examples: Eigen Faces
  - variant for classification: Linear Discriminant Analysis

